

# INFSCI 2595

Fall 2019

Information Sciences Building: Room 403

Lecture 03

Last week, we introduced the Bernoulli distribution

$$p(x|\mu) = \text{Bernoulli}(x|\mu) = \mu^x(1 - \mu)^{1-x}$$

- $x$  is a **binary variable**,  $x \in \{0, 1\}$
- $\mu$  is a probability and so is bounded:  $0 \leq \mu \leq 1$

We stepped through the Maximum Likelihood Estimate (MLE) of  $\mu$  given observations

- **N independent** observations,  $\mathbf{x} = \{x_1, x_2, \dots, x_n, \dots, x_N\}$
- We observe  $x = 1$  a total of  $M$  times.
- The MLE for the probability of the event is:

$$\mu_{ML} = \frac{M}{N}$$

But, let's ask a different question...

- Instead of asking, what's the probability  $x = 1$  (the EVENT)...
- Let's ask, what's the probability the event occurs a **specific number of times out of a specific number of trials?**

In terms of our college football example from last week...

- What's the probability of finding **exactly 1 Pitt fan out of 4 people?**

# Did we calculate this probability last week?

- Based on the following independent observations:

Person	Fan	$x$
1	PSU	0
2	PSU	0
3	Pitt	1
4	PSU	0

# Did we calculate this probability last week?

- Based on the following independent observations:

Person	Fan	$x$	$p(x \mu)$
1	PSU	0	$(1 - \mu)$
2	PSU	0	$(1 - \mu)$
3	Pitt	1	$\mu$
4	PSU	0	$(1 - \mu)$

# Did we calculate this probability last week?

- Based on the following independent observations:

Person	Fan	$x$	$p(x \mu)$
1	PSU	0	$(1 - \mu)$
2	PSU	0	$(1 - \mu)$
3	Pitt	1	$\mu$
4	PSU	0	$(1 - \mu)$

$$p(\mathbf{x}|\mu) = (1 - \mu)(1 - \mu)\mu(1 - \mu)$$



# Did we calculate this probability last week?

- Based on the following independent observations:

Person	Fan	$x$	$p(x \mu)$
--------	-----	-----	------------

Wait...is this the only way to observe 1 Pitt fan out of 4 people?

4	PSU	0	$(1 - \mu)$
---	-----	---	-------------

$$p(\mathbf{x}|\mu) = (1 - \mu)(1 - \mu)\mu(1 - \mu)$$

No! Multiple potential sequences of 4 people consist of exactly 1 Pitt fan.

Person 1	Person 2	Person 3	Person 4
Pitt	PSU	PSU	PSU
PSU	Pitt	PSU	PSU
PSU	PSU	Pitt	PSU
PSU	PSU	PSU	Pitt

Rewrite each of the potential sequences in terms of the encoded variable  $x$

$x_1$	$x_2$	$x_3$	$x_4$
1	0	0	0
0	1	0	0
0	0	1	0
0	0	0	1

Calculate the probability of each potential sequence assuming independent observations

$p(x_1 \mu)$	$p(x_2 \mu)$	$p(x_3 \mu)$	$p(x_4 \mu)$
$\mu$	$(1 - \mu)$	$(1 - \mu)$	$(1 - \mu)$
$(1 - \mu)$	$\mu$	$(1 - \mu)$	$(1 - \mu)$
$(1 - \mu)$	$(1 - \mu)$	$\mu$	$(1 - \mu)$
$(1 - \mu)$	$(1 - \mu)$	$(1 - \mu)$	$\mu$

Each of the potential sequences have the same probability!

$p(\mathbf{x} \mu)$
$\mu \cdot (1 - \mu)^3$
$\mu \cdot (1 - \mu)^3$
$\mu \cdot (1 - \mu)^3$
$\mu \cdot (1 - \mu)^3$

The probability of observing exactly 1 Pitt fan out of 4 people:

- Sum together the probabilities of each **potential** sequence:

$$4 \cdot \mu \cdot (1 - \mu)^3$$

- Next, what's the probability of finding exactly 2 Pitt fans out of 4 people?

List all potential sequences with 2 Pitt fans

$x_1$	$x_2$	$x_3$	$x_4$
1	1	0	0
0	1	1	0
0	0	1	1
1	0	1	0
0	1	0	1
1	0	0	1



Calculate the probability of each potential sequence assuming independent observations

$x_1$	$x_2$	$x_3$	$x_4$
$\mu$	$\mu$	$(1 - \mu)$	$(1 - \mu)$
$(1 - \mu)$	$\mu$	$\mu$	$(1 - \mu)$
$(1 - \mu)$	$(1 - \mu)$	$\mu$	$\mu$
$\mu$	$(1 - \mu)$	$\mu$	$(1 - \mu)$
$(1 - \mu)$	$\mu$	$(1 - \mu)$	$\mu$
$\mu$	$(1 - \mu)$	$(1 - \mu)$	$\mu$

Calculate the probability of each potential sequence assuming independent observations

$x_1$	$x_2$	$x_3$	$x_4$
	$\mu^2(1 - \mu)^2$		
	$\mu^2(1 - \mu)^2$		
	$\mu^2(1 - \mu)^2$		
	$\mu^2(1 - \mu)^2$		
	$\mu^2(1 - \mu)^2$		
	$\mu^2(1 - \mu)^2$		

The probability of observing exactly 2 Pitt fans out of 4 people:

- Sum together the probabilities of each **potential** sequence:

$$6 \cdot \mu^2 \cdot (1 - \mu)^2$$

How many potential sequences exist?

- Assume 4 people (trials).
- A person can be either a Pitt fan or a PSU fan (binary outcome).

$$2^4 = 16$$

Sequence ID	$x_1$	$x_2$	$x_3$	$x_4$
1	0	0	0	0
2	1	0	0	0
3	0	1	0	0
4	0	0	1	0
5	0	0	0	1
6	1	1	0	0
7	0	1	1	0
8	0	0	1	1
9	1	0	1	0
10	0	1	0	1
11	1	0	0	1
12	1	1	1	0
13	0	1	1	1
14	1	1	0	1
15	1	0	1	1
16	1	1	1	1

Sequence ID	$x_1$	$x_2$	$x_3$	$x_4$	Times $x = 1$
1	0	0	0	0	0
2	1	0	0	0	1
3	0	1	0	0	
4	0	0	1	0	
5	0	0	0	1	
6	1	1	0	0	2
7	0	1	1	0	
8	0	0	1	1	
9	1	0	1	0	
10	0	1	0	1	
11	1	0	0	1	
12	1	1	1	0	3
13	0	1	1	1	
14	1	1	0	1	
15	1	0	1	1	
16	1	1	1	1	4

Calculate the probability of observing  $x = 1$  exactly 0, 1, 2, 3, and 4 times.

Times $x = 1$	$p(\mathbf{x} \mu)$
0	$1 \cdot \mu^0 \cdot (1 - \mu)^4$
1	$4 \cdot \mu^1 \cdot (1 - \mu)^3$
2	$6 \cdot \mu^2 \cdot (1 - \mu)^2$
3	$4 \cdot \mu^3 \cdot (1 - \mu)^1$
4	$1 \cdot \mu^4 \cdot (1 - \mu)^0$

# WHAT PATTERNS DO YOU SEE??

Times $x = 1$	$p(\mathbf{x} \mu)$
0	$1 \cdot \mu^0 \cdot (1 - \mu)^4$
1	$4 \cdot \mu^1 \cdot (1 - \mu)^3$
2	$6 \cdot \mu^2 \cdot (1 - \mu)^2$
3	$4 \cdot \mu^3 \cdot (1 - \mu)^1$
4	$1 \cdot \mu^4 \cdot (1 - \mu)^0$



# WHAT PATTERNS DO YOU SEE??

The exponent on  $\mu$  equals the number of times  $x = 1$ .

The number of times  $x = 1$ , corresponds to the number of times we observed the EVENT.

Define the number of EVENTS to be  $m$ .

Times $x = 1$	$p(x \mu)$
0	$1 \cdot \mu^0 \cdot (1 - \mu)^4$
1	$4 \cdot \mu^1 \cdot (1 - \mu)^3$
2	$6 \cdot \mu^2 \cdot (1 - \mu)^2$
3	$4 \cdot \mu^3 \cdot (1 - \mu)^1$
4	$1 \cdot \mu^4 \cdot (1 - \mu)^0$

# WHAT PATTERNS DO YOU SEE??

The exponent on  $(1 - \mu)$  equals the number of TRIALS minus the number of EVENTS.

Corresponds to the number of times we did not observe the EVENT.

Define as  $N - m$ .

$m$	$p(\mathbf{x} \mu)$
0	$1 \cdot \mu^m \cdot (1 - \mu)^4$
1	$4 \cdot \mu^m \cdot (1 - \mu)^3$
2	$6 \cdot \mu^m \cdot (1 - \mu)^2$
3	$4 \cdot \mu^m \cdot (1 - \mu)^1$
4	$1 \cdot \mu^m \cdot (1 - \mu)^0$

# WHAT PATTERNS DO YOU SEE??

What about the coefficient out front?

Rewrite using:

$$\binom{4}{0} = 1, \binom{4}{1} = 4, \binom{4}{2} = 6$$
$$\binom{4}{3} = 4, \binom{4}{4} = 1$$

$m$	$p(\mathbf{x} \mu)$
0	$1 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
1	$4 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
2	$6 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
3	$4 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
4	$1 \cdot \mu^m \cdot (1 - \mu)^{N-m}$

# WHAT PATTERNS DO YOU SEE??

What about the coefficient out front?

Rewrite using:

$$\binom{4}{0} = 1, \binom{4}{1} = 4, \binom{4}{2} = 6$$
$$\binom{4}{3} = 4, \binom{4}{4} = 1$$

Which can be generalized using:

$$\binom{N}{m}$$

$m$	$p(\mathbf{x} \mu)$
0	$1 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
1	$4 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
2	$6 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
3	$4 \cdot \mu^m \cdot (1 - \mu)^{N-m}$
4	$1 \cdot \mu^m \cdot (1 - \mu)^{N-m}$

The probability distribution of  $m$  events out of  $N$  trials, given event probability  $\mu$ :

$$p(m|N, \mu) = \binom{N}{m} \mu^m (1 - \mu)^{N-m}$$


$$m \in \{0, \dots, N\}$$

Known as the Binomial distribution!

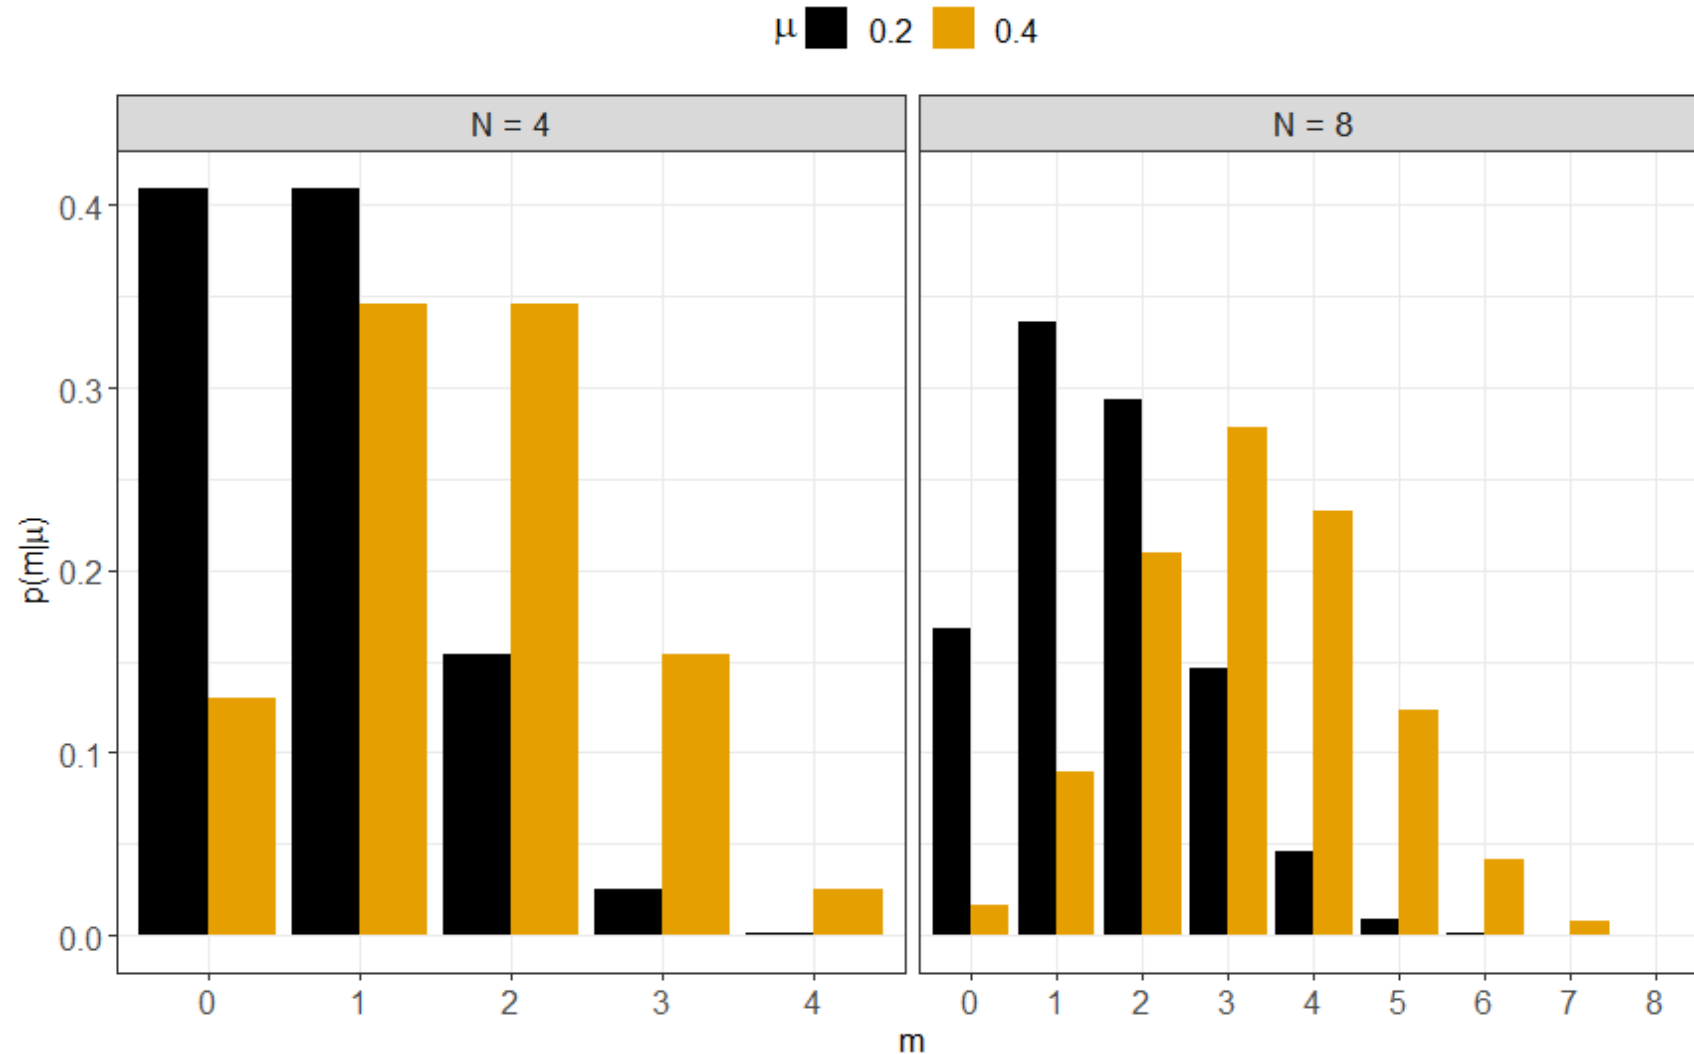
We derived the Binomial distribution starting from Bernoulli observations

- The Binomial distribution is a sequence of INDEPENDENT Bernoulli trials.
- We recover the Bernoulli distribution with  $N = 1$ . Thus,  $m = \{0,1\}$ .
- The Bernoulli is therefore a special case of the Binomial distribution.

# Binomial distribution for $N = 8$ and $\mu = 0.2$




# Binomial distribution two different $N$ 's and two different $\mu$ 's






Back to our college football example, remember that the game is located at State College...

- If we ask 4 people, and assume the **TRUE** probability of a Pitt fan is  $\mu = 0.2$ ...
- The probability of finding 0 Pitt fans is  $\approx 40\%$ !
- The probability of finding 2 Pitt fans is small but not negligible at  $\approx 15\%$ .



Back to our college football example, remember that the game is located at State College...

- If 0 out of 4 people are Pitt fans, our MLE for the probability would be  $\mu_{ML} = 0$ .
- If 2 out of 4 people are Pitt fans, our MLE for the probability would be  $\mu_{ML} = 0.5$ .
- Both estimates are not unrepresentative of  $\mu_{TRUE} = 0.2$ !



Our MLE is unreliable in this small data situation!

- How can we overcome this limitation?

Our MLE is unreliable in this small data situation!

- How can we overcome this limitation?
- Ask more people (collect more data)...but what if we cannot do that?

Our MLE is unreliable in this small data situation!

- How can we overcome this limitation?
- Ask more people (collect more data)...but what if we cannot do that?
- Could we make use of additional information?

Remember, the game is a **home** game for Penn State

- Thus, it is safe to anticipate more PSU fans than Pitt fans to be present at the game.
- How can we make use of this information in our analysis?

Remember, the game is a **home** game for Penn State

- Thus, it is safe to anticipate more PSU fans than Pitt fans to be present at the game.
- How can we make use of this information in our analysis?

**Bayesian statistics!**

# Bayesian formulation for estimating $\mu$

- We want to update our prior belief about  $\mu$  based on observations.

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$



# Bayesian formulation for estimating $\mu$

- We want to update our prior belief about  $\mu$  based on observations.

Posterior  $\propto$  Likelihood  $\times$  Prior

Based on the binomial distribution as the likelihood

$$p(\mu|m, N) \propto \text{Binomial}(m|N, \mu)p(\mu)$$

# Bayesian formulation for estimating $\mu$

- We want to update our prior belief about  $\mu$  based on observations.

Posterior  $\propto$  Likelihood  $\times$  Prior

Based on the binomial distribution as the likelihood

$$p(\mu|m, N) \propto \text{Binomial}(m|N, \mu)p(\mu)$$

Or, based on independent Bernoulli trials as the likelihood

$$p(\mu|\mathbf{x}) \propto \prod_{n=1}^N \{\text{Bernoulli}(x_n|\mu)\} p(\mu)$$

# Bayesian formulation for estimating $\mu$

- We want to update our prior belief about  $\mu$  based on observations.

Posterior  $\propto$  Likelihood  $\times$  Prior

$$p(\mu|m, N) \propto \text{Binomial}(m|N, \mu)p(\mu)$$

Use this  
formulation  
now.

$$p(\mu|\mathbf{x}) \propto \prod_{n=1}^N \{\text{Bernoulli}(x_n|\mu)\} p(\mu)$$

# Bayesian formulation for estimating $\mu$

- We want to update our prior belief about  $\mu$  based on observations.

Posterior  $\propto$  Likelihood  $\times$  Prior

$$p(\mu|m, N) \propto \text{Binomial}(m|N, \mu) p(\mu)$$

We know how to write out the likelihood...but what about the prior,  $p(\mu)$ ?

How can we specify a prior belief about  $\mu$ ?

We will use a BETA distribution to encode our **PRIOR** belief on the probability,  $\mu$

# Beta distribution

- The beta distribution is a probability density function (pdf) for continuous variables **BOUNDED** between 0 and 1.

# Beta distribution


- The beta distribution is a probability density function (pdf) for continuous variables **BOUNDED** between 0 and 1.
- It is a flexible distribution capable of a wide variety of shapes.

# Beta distribution


- The beta distribution is a probability density function (pdf) for continuous variables **BOUNDED** between 0 and 1.
- It is a flexible distribution capable of a wide variety of shapes.
- The shape is controlled by the hyperparameters  $\alpha$  and  $\beta$ .
  - The Bishop book denotes these two parameters as  $a$  and  $b$ .



# Example shapes of the beta distribution



# Example shapes of the beta distribution



The beta pdf...

$$p(\mu|a, b) = \text{Beta}(\mu|a, b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \mu^{a-1} (1-\mu)^{b-1}$$

The beta pdf...looks rather familiar...

$$p(\mu|a, b) = \text{Beta}(\mu|a, b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \mu^{a-1} (1-\mu)^{b-1}$$

- Focus on the terms involving  $\mu$ :

$$\text{Beta}(\mu|a, b) \propto \mu^{a-1} (1-\mu)^{b-1}$$

The beta distribution has the same functional form as the Binomial distribution!

$$\text{Beta}(\mu|a, b) \propto \mu^{a-1}(1 - \mu)^{b-1}$$



$$\text{Binomial}(m|\mu, N) \propto \mu^m(1 - \mu)^{N-m}$$

- The beta distribution is the **conjugate prior** of the binomial likelihood.

The beta distribution has the same functional form as the Binomial distribution!

$$\text{Beta}(\mu|a, b) \propto \mu^{a-1}(1 - \mu)^{b-1}$$



$$\text{Binomial}(m|\mu, N) \propto \mu^m(1 - \mu)^{N-m}$$

- A conjugate prior is useful because the resulting posterior distribution will have the same functional form as the prior.

The beta distribution has the same functional form as the Binomial distribution!

$$\text{Beta}(\mu|a, b) \propto \mu^{a-1}(1 - \mu)^{b-1}$$



$$\text{Binomial}(m|\mu, N) \propto \mu^m(1 - \mu)^{N-m}$$

- The posterior will therefore also be a **beta distribution!**


Posterior distribution on  $\mu$

$$p(\mu|m, N) = \text{Beta}(\mu|a + m, b + (N - m))$$



Posterior distribution on  $\mu$

$$p(\mu|m, N) = \text{Beta}(\mu | \underbrace{a + m}_{a_{new}}, \underbrace{b + (N - m)}_{b_{new}})$$



$$p(\mu|m, N) = \text{Beta}(\mu | a_{new}, b_{new})$$

# Beta distribution hyperparameter interpretations

- $a$  is added to the number of Pitt fans,  $m$ , or more generally the number of observed EVENTS.
- $b$  is added to the number of PSU fans,  $N - m$ , or more generally the number of times we did NOT observe the EVENT.

# Beta distribution hyperparameter interpretations

- $a$  is added to the number of Pitt fans,  $m$ , or more generally the number of observed EVENTS.
  - $a$  is therefore the *a priori* number of EVENTS!!
- $b$  is added to the number of PSU fans,  $N - m$ , or more generally the number of times we did NOT observe the EVENT.
  - $b$  is therefore the *a priori* number of NON-EVENTS!!

We could have reached the same interpretations by considering the mean...

- The expected value (mean) of the Beta distribution is:

$$\mathbb{E}[\mu|a, b] = \frac{a}{a + b}$$

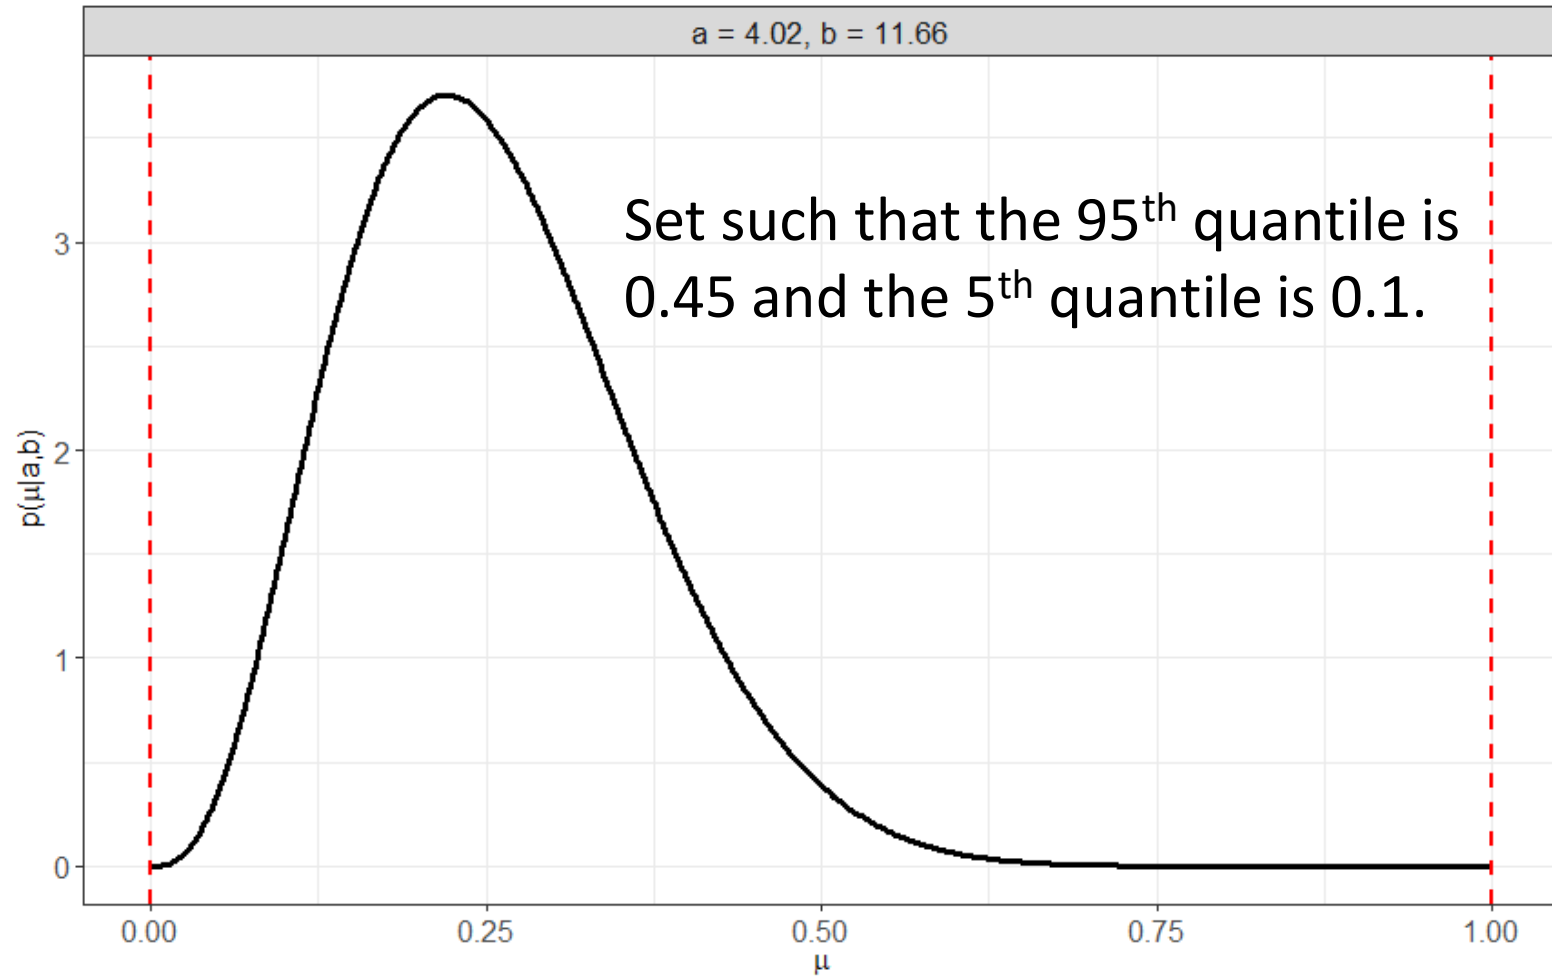
We could have reached the same interpretations by considering the mean...

- The expected value (mean) of the Beta distribution is:

$$\mathbb{E}[\mu|a, b] = \frac{a}{a + b} \Rightarrow \frac{\text{Number of events!}}{\text{Number of trials!}}$$

$b$  is therefore the number of  
NON-EVENTS, or times  $x = 0$ !

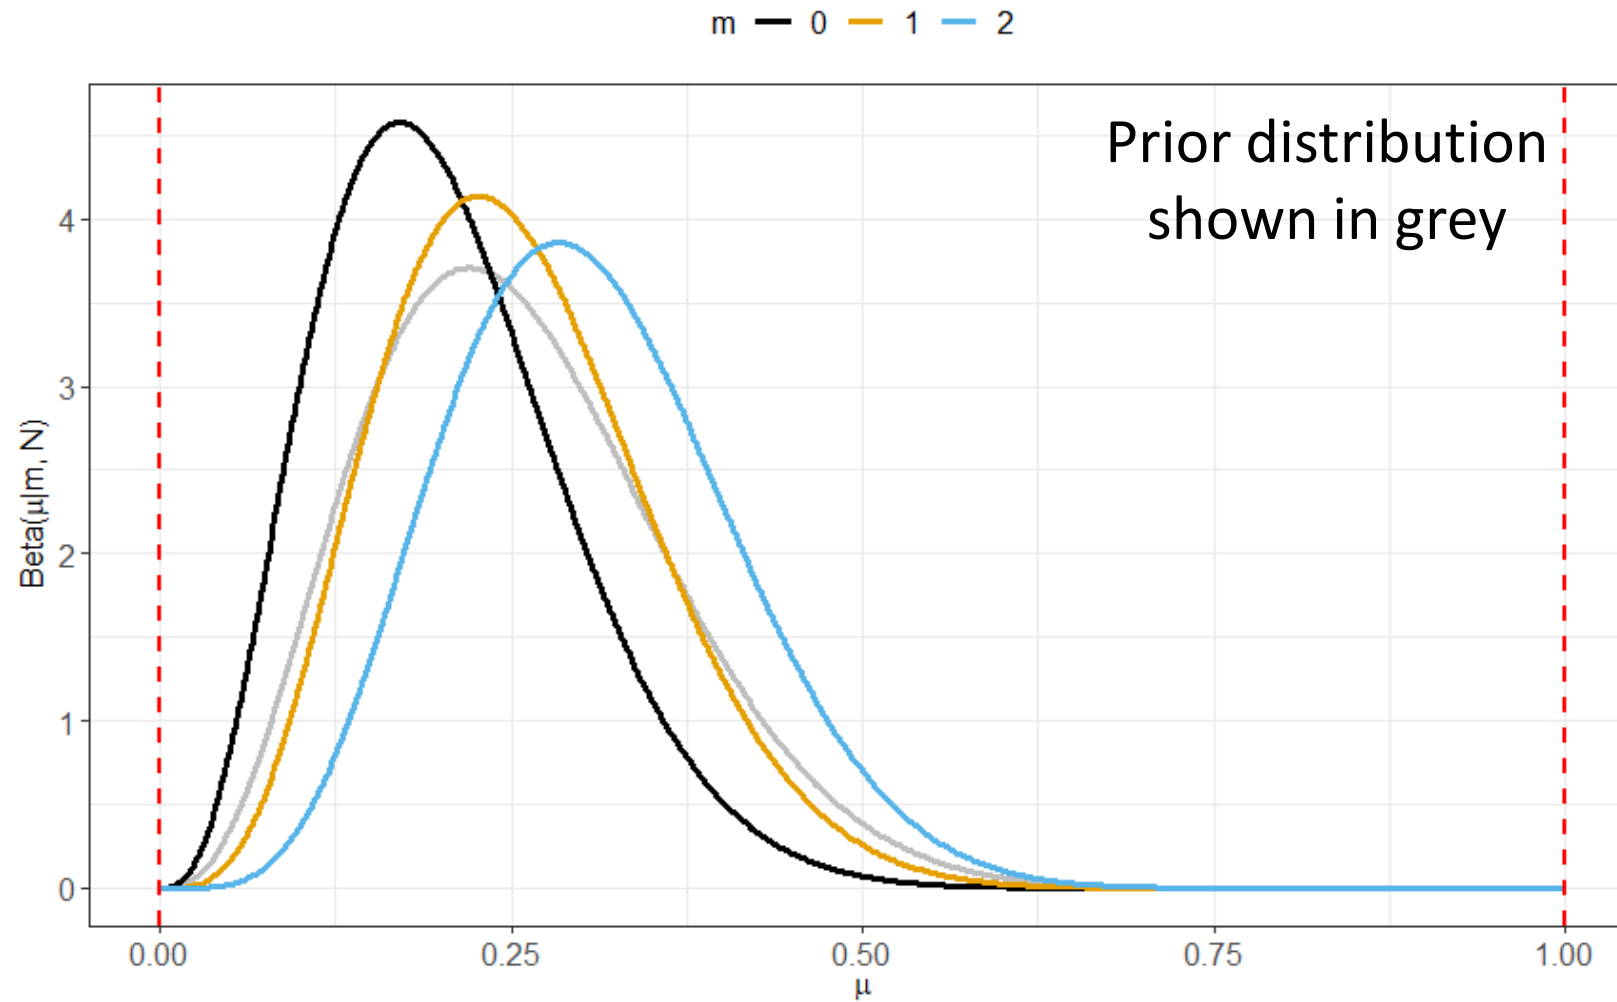
Set our prior such that we feel the probability of finding a Pitt fan is greater than 0 but less than 0.5



We will update our belief about  $\mu$  under three different circumstances

- As we saw, the posterior distribution on  $\mu$  given the observations is a Beta distribution.
- Let's compare the resulting Beta distributions based on observing  $m = 0, 1$ , and  $2$ .
- Thus, what's our **updated belief** if we found 0 Pitt fans, vs 1 Pitt fan, vs 2 Pitt fans.

$\mu$  posterior distribution given  $m$  and  $N = 4$







# Summarize the Beta distributions

- Calculate summary statistics for each Beta distribution.
- Represent uncertainty with **credible intervals**:
  - Middle 50% interval – spans the 25<sup>th</sup> through 75<sup>th</sup> quantiles
  - Middle 90% interval – spans the 5<sup>th</sup> through 95<sup>th</sup> quantiles
- Represent the central tendency two ways:
  - Median – the 50<sup>th</sup> quantile
  - Mean (average value)


# Visualize the Beta distribution summary statistics




# Visualize the Beta distribution summary statistics



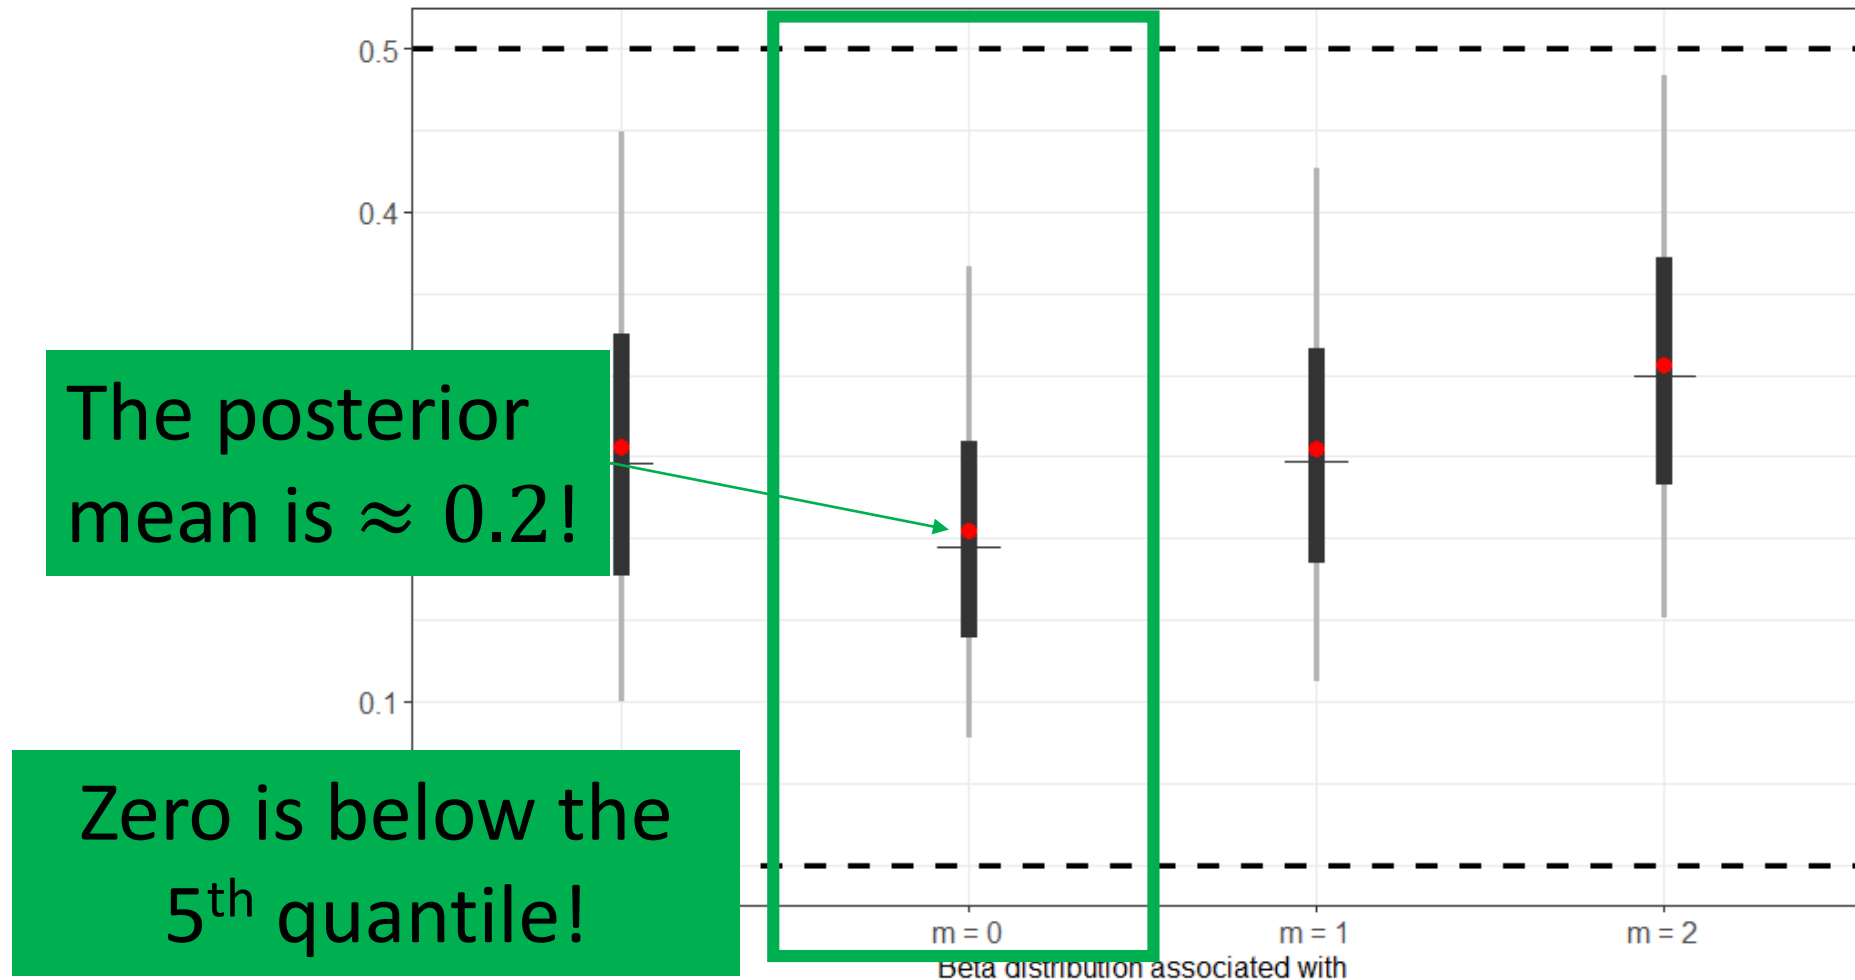
# Zoom in



*A priori* we believe the mean is  $\approx 0.25$




If we observed  $m = 0$  out of  $N = 4$



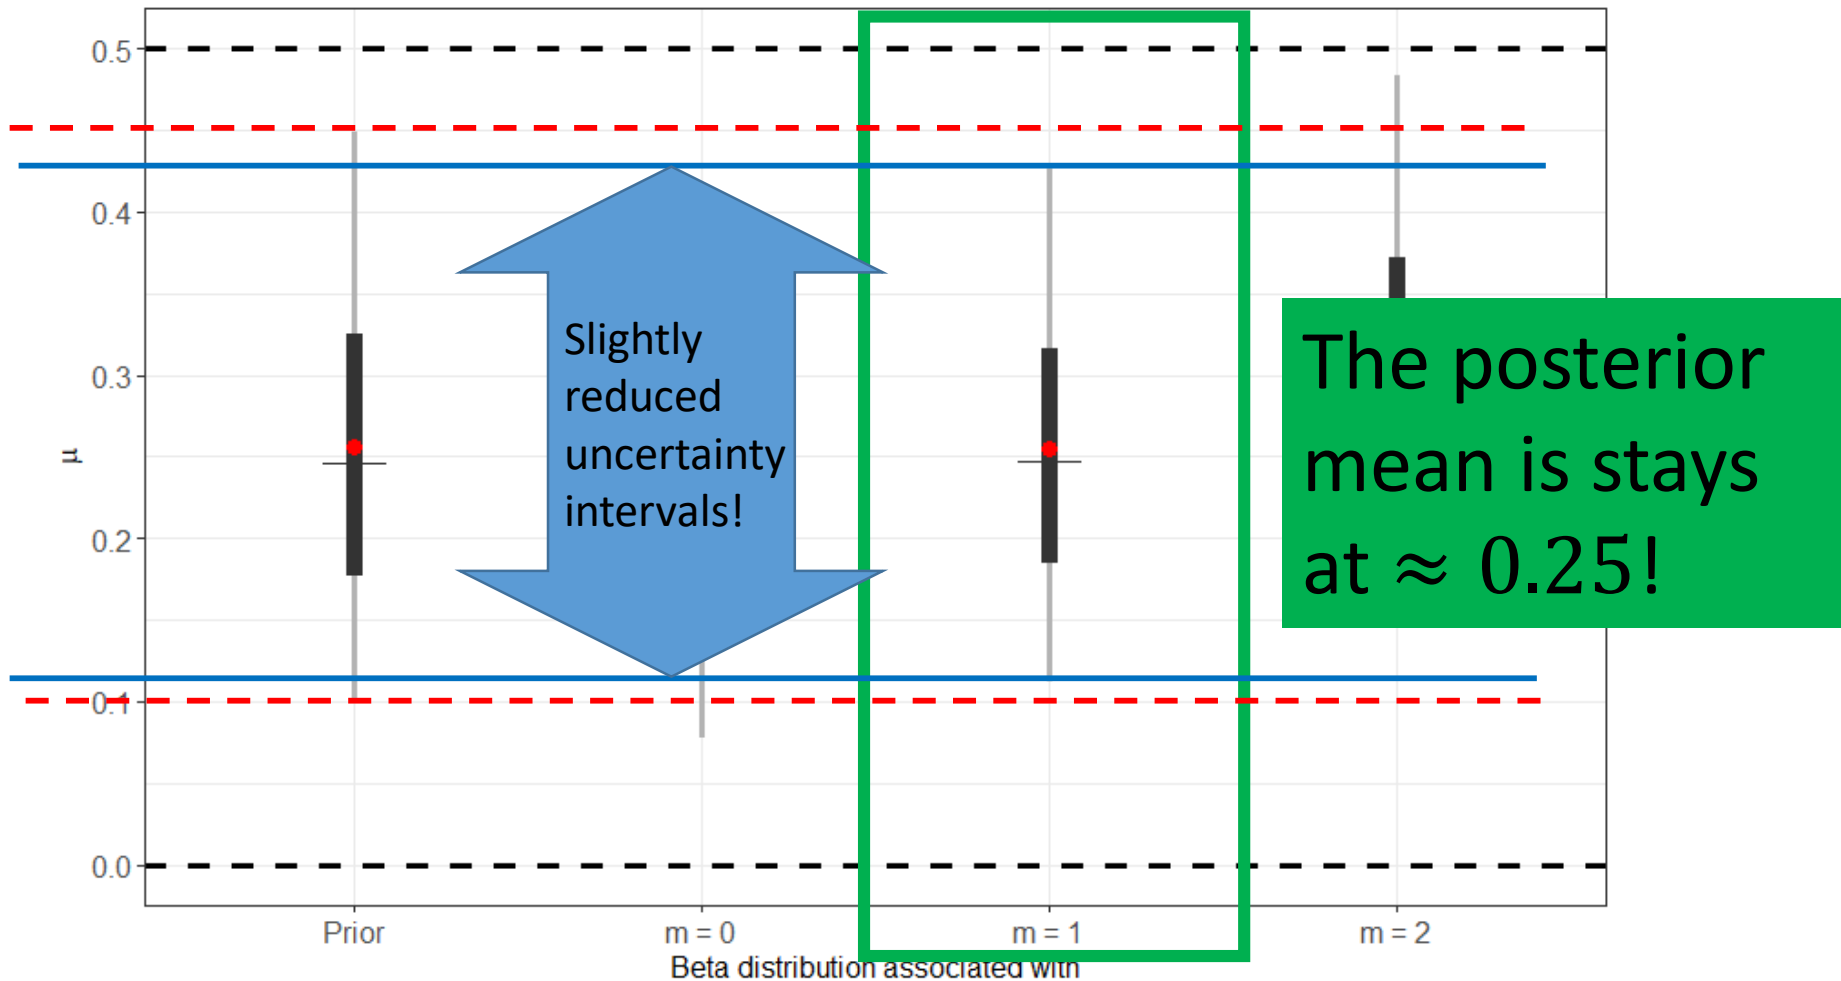
If we observed  $m = 2$  out of  $N = 4$

The prior keeps 0.5 above the 95<sup>th</sup> quantile!

The posterior mean is  $\approx 0.3$ !



If we observed  $m = 1$  out of  $N = 4$





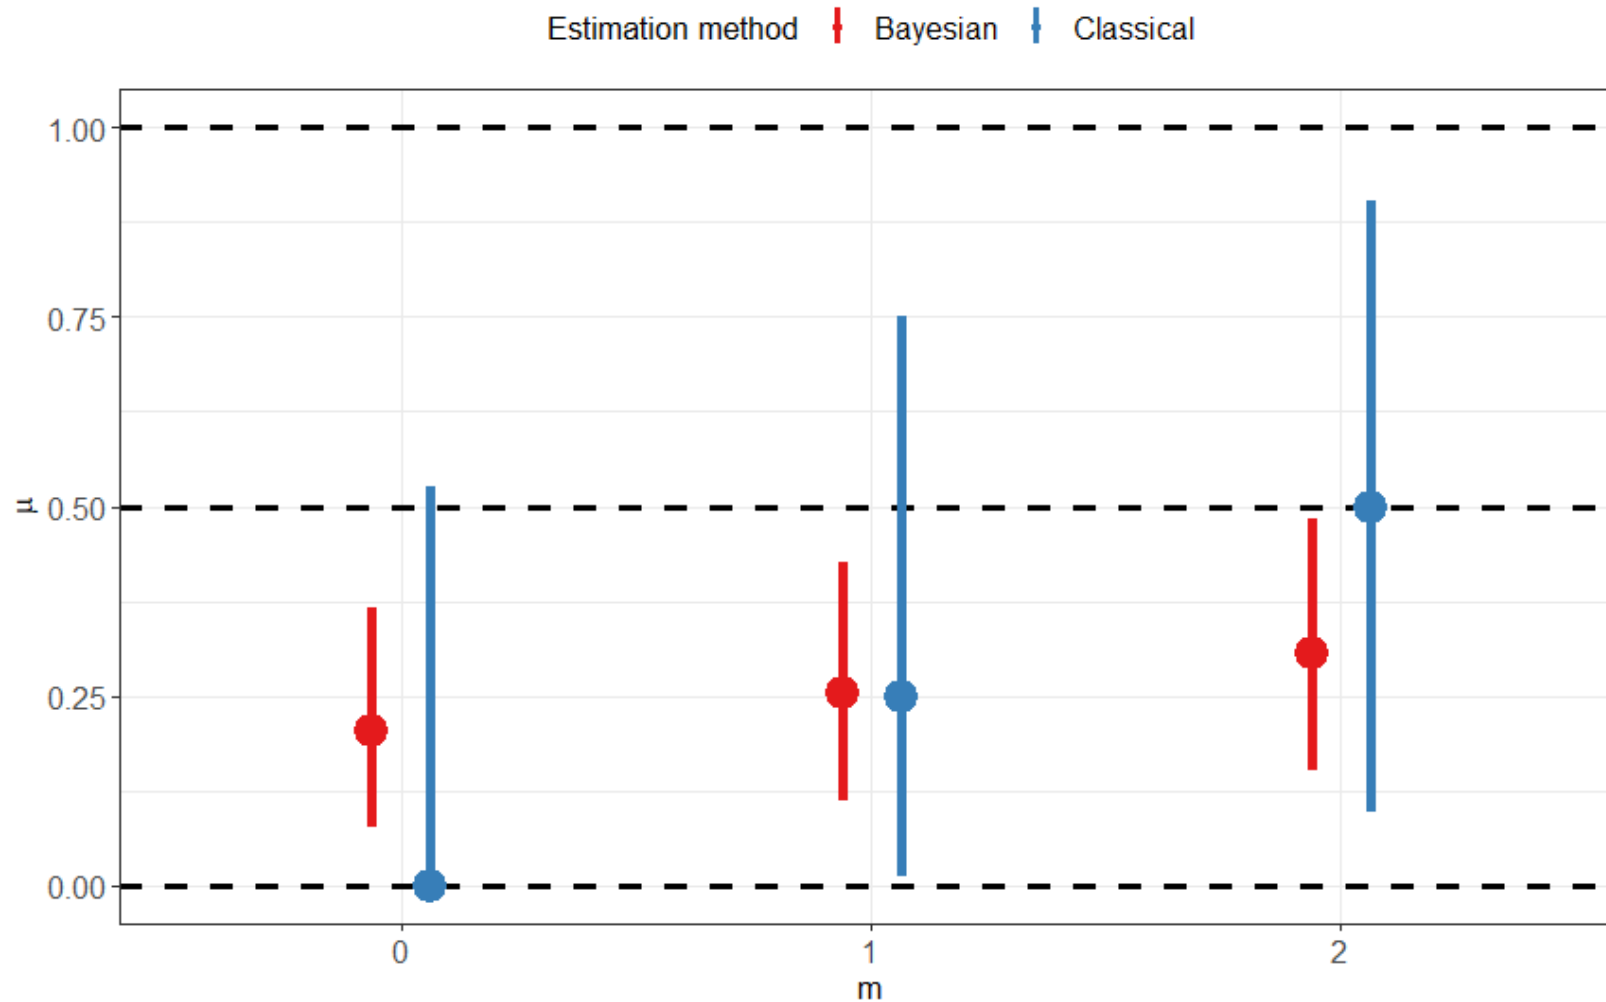
# We introduced discussing uncertainty from a Bayesian framework

- However, classical or frequentist statistics also have ways for estimating uncertainty.
- Uncertainty usually represented by **confidence intervals**.
- How do 90% confidence intervals compare with the *posterior credible intervals* in our college football example?

# Confidence interval calculation

- The 90% confidence intervals are calculating using the Clopper-Pearson method, through R's `binom.test()` function.
- Please see `?binom.test` for more discussion around the method.

# 90% credible intervals (red) compared with the 90% confidence intervals (blue)



# 90% credible intervals (red) compared with the 90% confidence intervals (blue)

