

INFSCI 2595

Fall 2019

Information Sciences Building: Room 403

Lecture 05

Last week...

- We discussed two continuous probability distributions
 - Beta distribution
 - Gaussian distribution
- Please see the [Probability Density Function Review](#) on the Course Github page to see how to work with pdf's in \mathbb{R} .

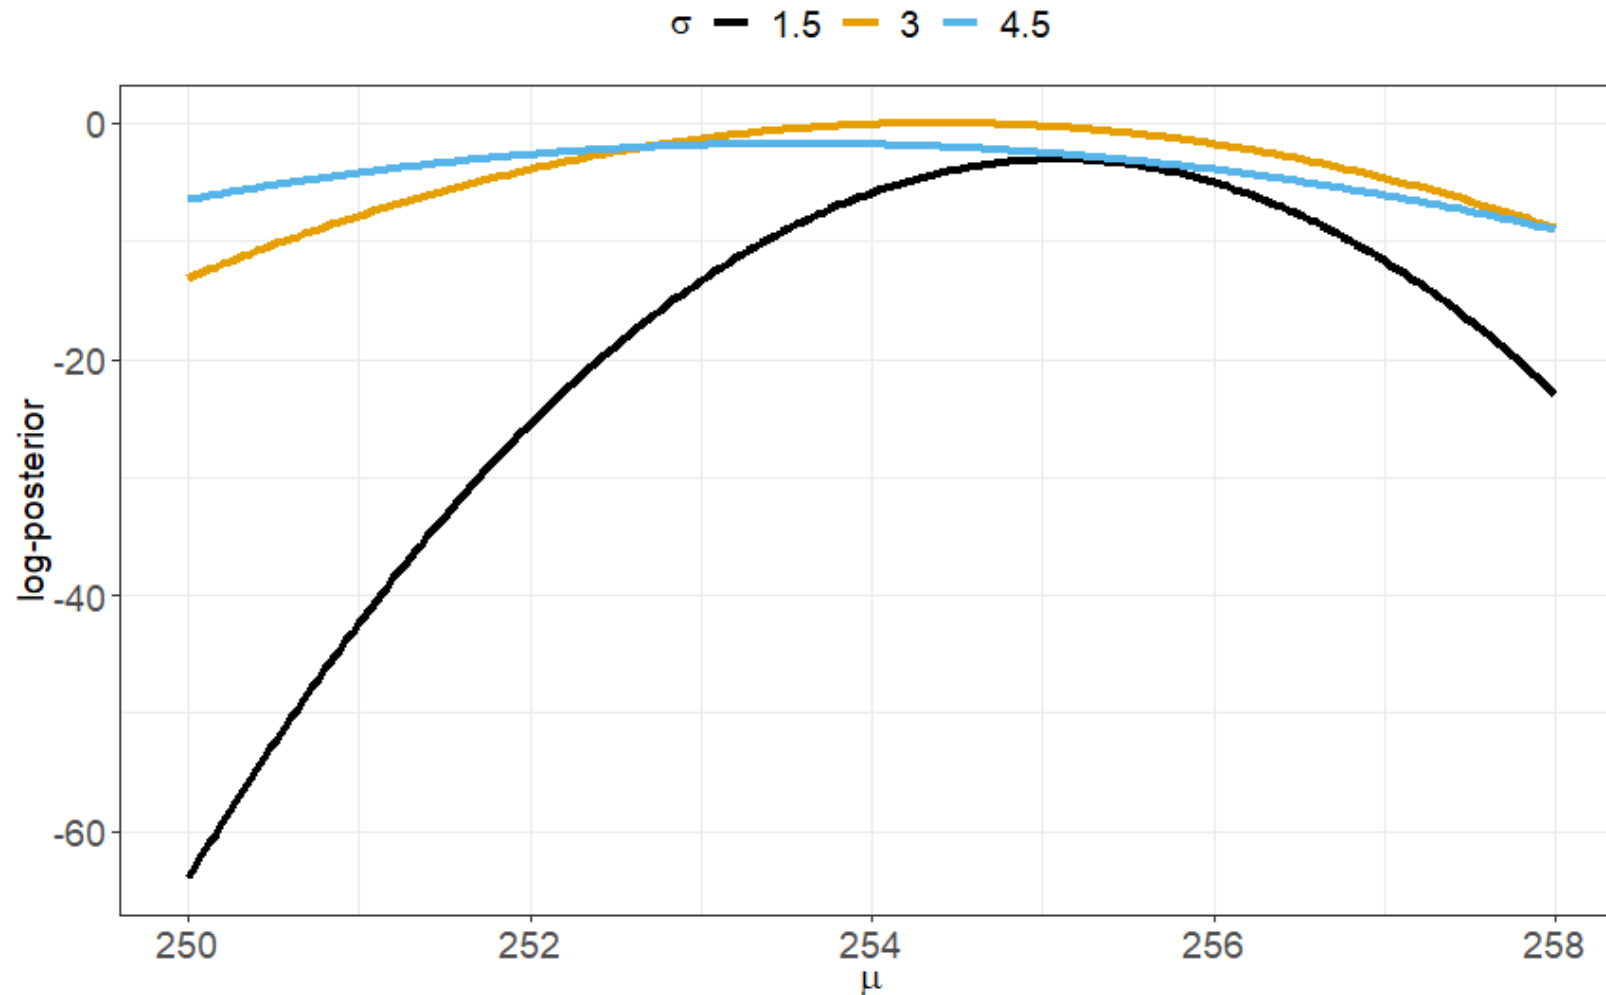
Last week...

- We discussed the normal model with unknown mean, μ , and known standard deviation, σ .
- Please see the [normal-normal model](#) supplemental material for more details and example code for evaluating log-posterior densities.

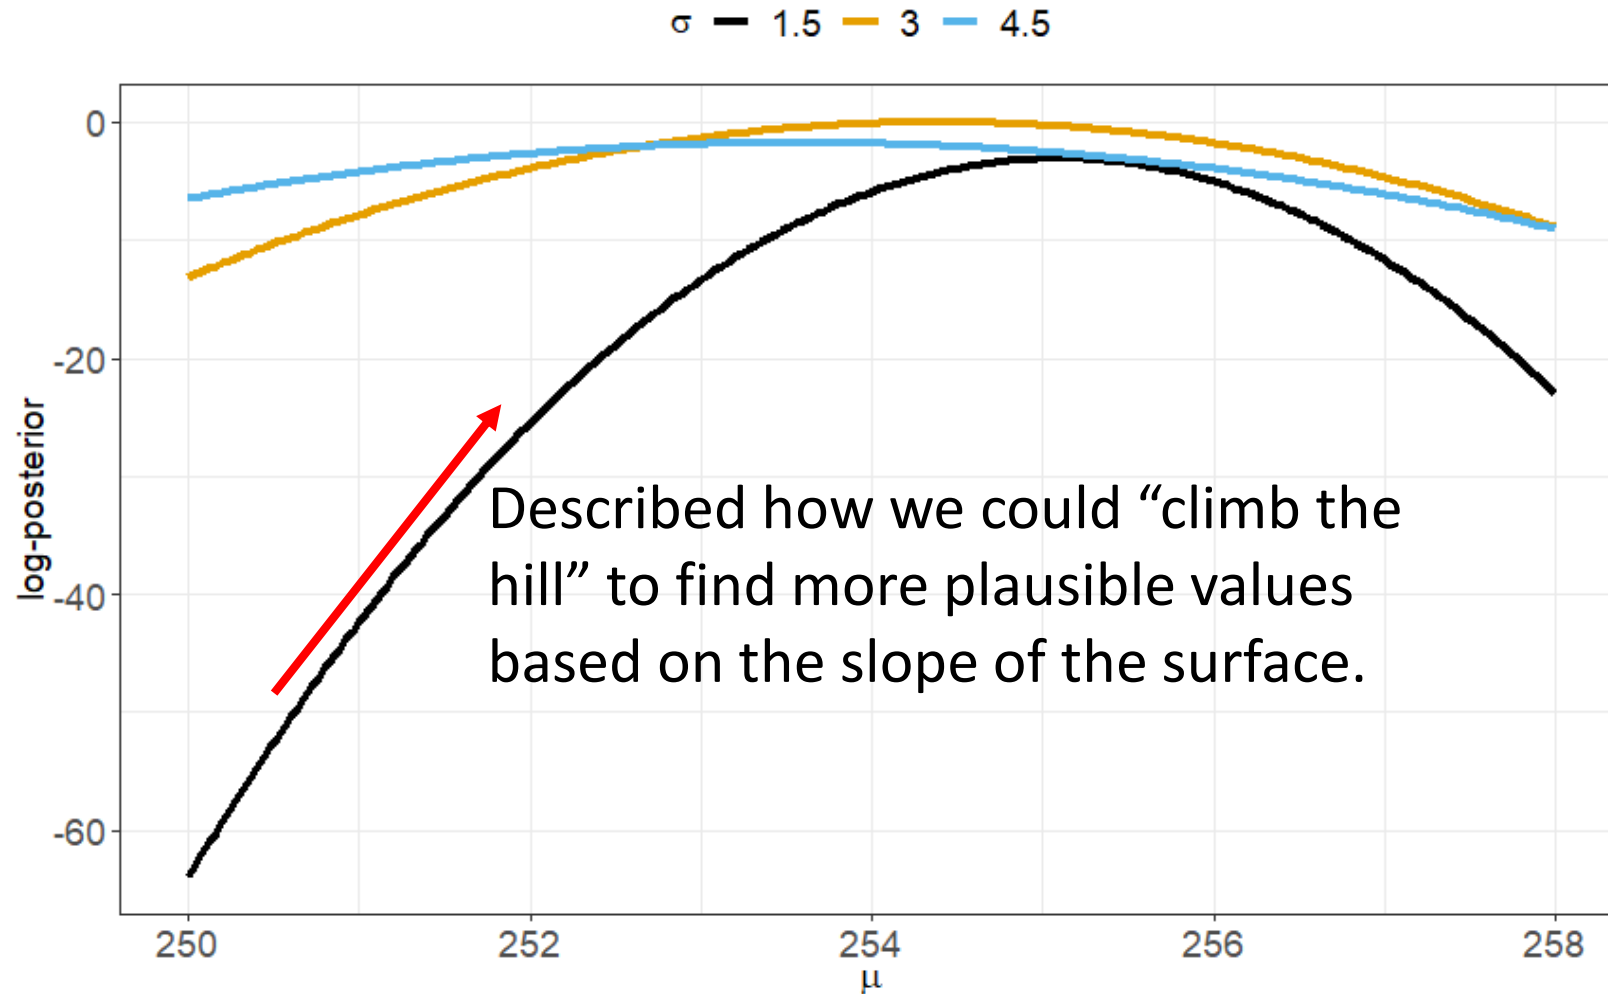
We concluded last week by introducing the grid approximation

- We used the **grid approximation**, or *direct sampling*, to estimate the joint posterior distribution for an unknown mean and unknown standard deviation.
- Allowed us to estimate the probability that I weigh less than 255 pounds based on 10 observations.

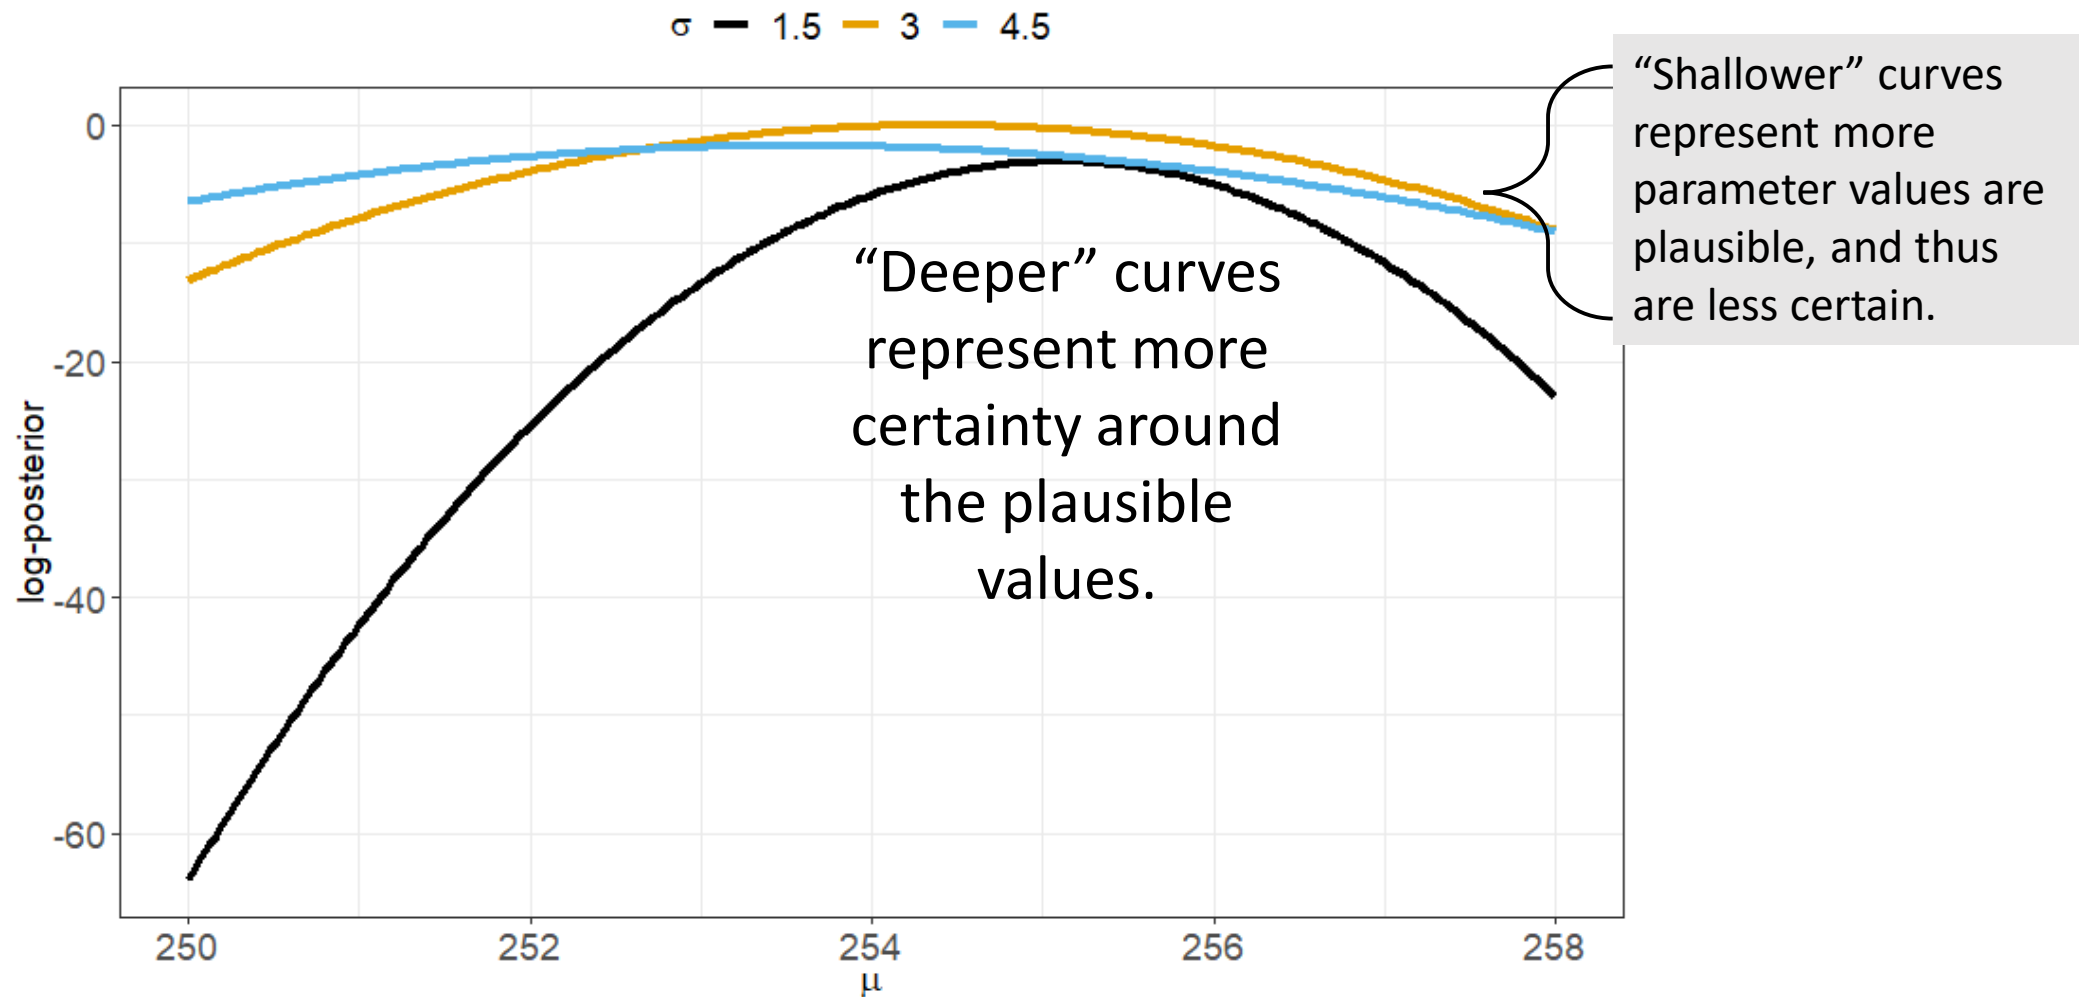
In our example, we visualized the log-posterior density with respect to the parameters



We discussed how higher log-posterior values correspond to *more probable* parameter values



We discussed how the shape or curvature of the surface represented the *certainty of plausible* parameter values



Today, we will introduce a method to quantify these aspects.

- This method will allow us to approximate the posterior distribution centered around the most probable value or *Max A Posteriori* (MAP) estimate.
- This method will approximate our uncertainty in the parameter around the MAP based on the log-posterior surface curvature.

Laplace, Quadratic, or Normal approximation

- Approximate the joint posterior distribution with a **Multivariate Normal** (MVN) centered on the **MAP**.
- Benefits:
 - Straightforward to implement.
 - Relatively fast to execute.
 - Scales to a moderate number of variables.
- Cons:
 - Let's see with an example later...

First things first...what's a MVN?

- Generalization of the Gaussian distribution to more than 1 dimension.
- Each dimension (variable) is a Gaussian and each subset of variables are MVN.

MVN density function

- There are D elements to the vector of variables:

$$\mathbf{x} = \{x_1, x_2, \dots, x_d, \dots, x_D\}$$

- **IMPORTANT:** D refers to the number of variables, NOT the number of observations!

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d, \boxed{\boldsymbol{\mu}} | \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Vector of means
associated with each
element in the x-vector:
 $\boldsymbol{\mu} = \{\mu_1, \mu_2, \dots, \mu_d, \dots, \mu_D\}$

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

$D \times D$ (variance-) covariance matrix between all elements of the \mathbf{x} -vector.

Off-diagonal elements store the covariance between the variables.

Main-diagonal elements store the variance of each element in the \mathbf{x} -vector

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D} |\boldsymbol{\Sigma}|} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Determinant of the
covariance matrix.

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Inverse of the
covariance matrix.

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Transpose of the
 $(\mathbf{x} - \boldsymbol{\mu})$ vector

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

What's this?

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

Multidimensional generalization of
the 1-D Gaussian term:

$$\left(\frac{x - \mu}{\sigma} \right)^2$$

Break down the terms in the density

$$p(x_1, x_2, \dots, x_d | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^D |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$

$\sqrt{(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})}$ is a generalized distance known as the Mahalanobis distance

Bivariate Gaussian – 2D case

- $D = 2$ the vector of elements becomes: $\mathbf{x} = \{x_1, x_2\}$
- Define the correlation coefficient between the two variables as, ρ .
- The mean vector and covariance matrix are:

$$\boldsymbol{\mu} = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}, \boldsymbol{\Sigma} = \begin{bmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{bmatrix}$$

Bivariate Gaussian – marginal distributions

- Each variable has a marginal Gaussian distribution:

$$x_1 | \mu_1, \sigma_1 \sim \text{normal}(x_1 | \mu_1, \sigma_1)$$

$$x_2 | \mu_2, \sigma_2 \sim \text{normal}(x_2 | \mu_2, \sigma_2)$$

Holds for higher dimensions!

Bivariate Gaussian – conditional distribution

- The conditional distribution of one variable given the other...is also a Gaussian!

$$x_1 | x_2, \boldsymbol{\mu}, \boldsymbol{\Sigma} \sim \mathcal{N} \left(\mu_1 + \frac{\sigma_1}{\sigma_2} \rho (x_2 - \mu_2), (1 - \rho^2) \sigma_1^2 \right)$$

Holds for higher dimensions!

How can we use a MVN to help us?

- Denote the parameters of interest as $\boldsymbol{\theta}$.
- For the weight example from last week, $\boldsymbol{\theta} = \{\mu, \sigma\}$.
- Assume we can find the posterior mode, or MAP, denote as $\hat{\boldsymbol{\theta}}$.

How can we use a MVN to help us?

- A *second-order Taylor series expansion* of the log-posterior around the posterior mode.

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{g}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

How can we use a MVN to help us?

- A *second-order Taylor series expansion* of the log-posterior around the posterior mode.

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{g}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

Gradient vector of the log-posterior with respect to each parameter of interest, evaluated at the posterior mode.

For the weight example: $\mathbf{g} = \left\{ \frac{\partial}{\partial \mu} (\log[p(\boldsymbol{\theta}|\mathbf{x})]), \frac{\partial}{\partial \sigma} (\log[p(\boldsymbol{\theta}|\mathbf{x})]) \right\}$

How can we use a MVN to help us?

- A *second-order Taylor series expansion* of the log-posterior around the posterior mode.

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{g}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

Gradient vector of the log-posterior with respect to each parameter of interest, evaluated at the posterior mode.

At $\boldsymbol{\theta} = \hat{\boldsymbol{\theta}}$ the gradient, \mathbf{g} , by definition, is equal to...

How can we use a MVN to help us?

- A *second-order Taylor series expansion* of the log-posterior around the posterior mode.

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{g}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

Gradient vector of the log-posterior with respect to each parameter of interest, evaluated at the posterior mode.

At $\boldsymbol{\theta} = \hat{\boldsymbol{\theta}}$ the gradient, \mathbf{g} , by definition, is equal to...**0!!!!!!!**

The second-order Taylor series expansion simplifies to:

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H} \Big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

The second-order Taylor series expansion simplifies to:

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H} \Big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

The **Hessian matrix** of the log-posterior with respect to the pair-wise parameter combinations evaluated at the posterior mode.

The Hessian matrix is a matrix of second derivatives and represents the **local curvature** of the surface.

The second-order Taylor series expansion simplifies to:

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H} \Big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

For the weight example, the Hessian matrix is a 2×2 matrix of second derivatives:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2}{\partial \mu^2} (\log[p(\boldsymbol{\theta}|\mathbf{x})]) & \frac{\partial^2}{\partial \mu \partial \sigma} (\log[p(\boldsymbol{\theta}|\mathbf{x})]) \\ \frac{\partial^2}{\partial \sigma \partial \mu} (\log[p(\boldsymbol{\theta}|\mathbf{x})]) & \frac{\partial^2}{\partial \sigma^2} (\log[p(\boldsymbol{\theta}|\mathbf{x})]) \end{bmatrix}$$

What does this expression remind us of...?

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \mathbf{H} \Big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

What does this expression remind us of...?

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \boxed{\hat{\boldsymbol{\theta}}})^T \mathbf{H} \Big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}} (\boldsymbol{\theta} - \boxed{\hat{\boldsymbol{\theta}}})$$

A vector of mean values!

What does this expression remind us of...?

$$\log[p(\boldsymbol{\theta}|\mathbf{x})] \approx \log[p(\hat{\boldsymbol{\theta}}|x)] + \frac{1}{2} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})^T \boxed{\mathbf{H} \big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}}} (\boldsymbol{\theta} - \hat{\boldsymbol{\theta}})$$

Negative inverse of the
covariance matrix!

The negative of the Hessian matrix has a special
name...the **observed information matrix**.

The Laplace approximation

$$p(\boldsymbol{\theta}|\mathbf{x}) \approx \mathcal{N}\left(\hat{\boldsymbol{\theta}}, \left[-\mathbf{H}\big|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}}\right]^{-1}\right)$$

Let's apply the Laplace approximation to our weight example from last week.

- The posterior distribution on μ, σ was proportional to:

$$p(\mu, \sigma | \mathbf{x}) \propto \prod_{n=1}^N \{\text{normal}(x_n | \mu, \sigma)\} \cdot \text{normal}(\mu | \mu_0, \tau_0) \cdot \text{uniform}(\sigma | l, u)$$

- Last week we used the following hyperparameters:

$$\mu_0 = 250, \tau_0 = 2$$

$$l = 1, u = 5$$

Let's apply the Laplace approximation to our weight example from last week.

- The posterior distribution on μ, σ was proportional to:

$$p(\mu, \sigma | \mathbf{x}) \propto \prod_{n=1}^N \{\text{normal}(x_n | \mu, \sigma)\} \cdot \text{normal}(\mu | \mu_0, \tau_0) \cdot \text{uniform}(\sigma | l, u)$$

- However, this week, increase the upper bound on σ :

$$\mu_0 = 250, \tau_0 = 2$$

$$l = 1, u = 20$$

Last week, we used 10 observations...

- However, this week we will use just the **first observation**, $N = 1$:

$$x_1 = 260.30$$

- Given this single observation, and our prior specification, **what is the posterior joint distribution on μ and σ ?**

Let's first visualize the log-posterior surface over a fine grid of μ and σ

- Thus, we will repeat the **grid approximation** from last week using our new assumptions and just the single observation.
- Define a function for evaluating the log-posterior.
- Define the following grid: $\mu \in [240, 260]$, $\sigma \in [1, 20]$

The log-posterior:

$$\log[p(\mu, \sigma | \mathbf{x})] \propto \sum_{n=1}^N \{\log[\text{normal}(x_n | \mu, \sigma)]\} + \log[\text{normal}(\mu | \mu_0, \tau_0)] + \log[\text{uniform}(\sigma | l, u)]$$

In R, define the function `my_logpost()`

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```


In R, define the function `my_logpost()`

The parameters, μ and σ , are defined as the first and second elements of the `theta` vector.

The second argument, `my_info`, is a list which stores all information required to evaluate the log-posterior.

- Contains the vector `xobs` which stores the observations.
- Contains of the hyperparameters to the prior on μ .
- Contains the hyperparameters to the prior on σ .

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

In R, define the function `my_logpost()`

$$\sum_{n=1}^N \{\log[\text{normal}(x_n | \mu, \sigma)]\}$$

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                        mean = lik_mu,
29                        sd = lik_sigma,
30                        log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                      mean = my_info$mu_0,
35                      sd = my_info$tau_0,
36                      log = TRUE) +
37     dunif(x = lik_sigma,
38           min = my_info$sigma_lwr,
39           max = my_info$sigma_upr,
40           log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

In R, define the function `my_logpost()`

$\log[\text{normal}(\mu|\mu_0, \tau_0)]$

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                        mean = lik_mu,
29                        sd = lik_sigma,
30                        log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                      mean = my_info$mu_0,
35                      sd = my_info$tau_0,
36                      log = TRUE) +
37     dunif(x = lik_sigma,
38           min = my_info$sigma_lwr,
39           max = my_info$sigma_upr,
40           log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

In R, define the function `my_logpost()`

$\log[\text{uniform}(\sigma|l, u)]$

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

Wrap `my_logpost ()` in a function to help manage its execution.

```
46 ### create a wrapper function which will allow evaluating the log-posterior  
47 ### over a defined grid of parameter values  
48 eval_logpost <- function(mu_val, sigma_val, my_info)  
49 {  
50   my_logpost(c(mu_val, sigma_val), my_info)  
51 }  
52
```

Create the full-factorial grid of parameter values with the `expand.grid()` function

```
68 ### define a grid of points to use
69 param_grid <- expand.grid(mu = seq(240, 260, length.out = 201),
70                           sigma = seq(info_use$sigma_lwr,
71                                       info_use$sigma_upr,
72                                       length.out = 201),
73                           KEEP.OUT.ATTRS = FALSE,
74                           stringsAsFactors = FALSE) %>%
75   as.data.frame() %>% tbl_df()
76
```


Define the hyperparameters and set the single observation appropriately

```
57 Nuse <- 1
58 xuse <- x[1:Nuse]
59
60 info_use <- list(
61   xobs = xuse,
62   mu_0 = 250,
63   tau_0 = 2,
64   sigma_lwr = 1,
65   sigma_upr = 20
66 )
67
```

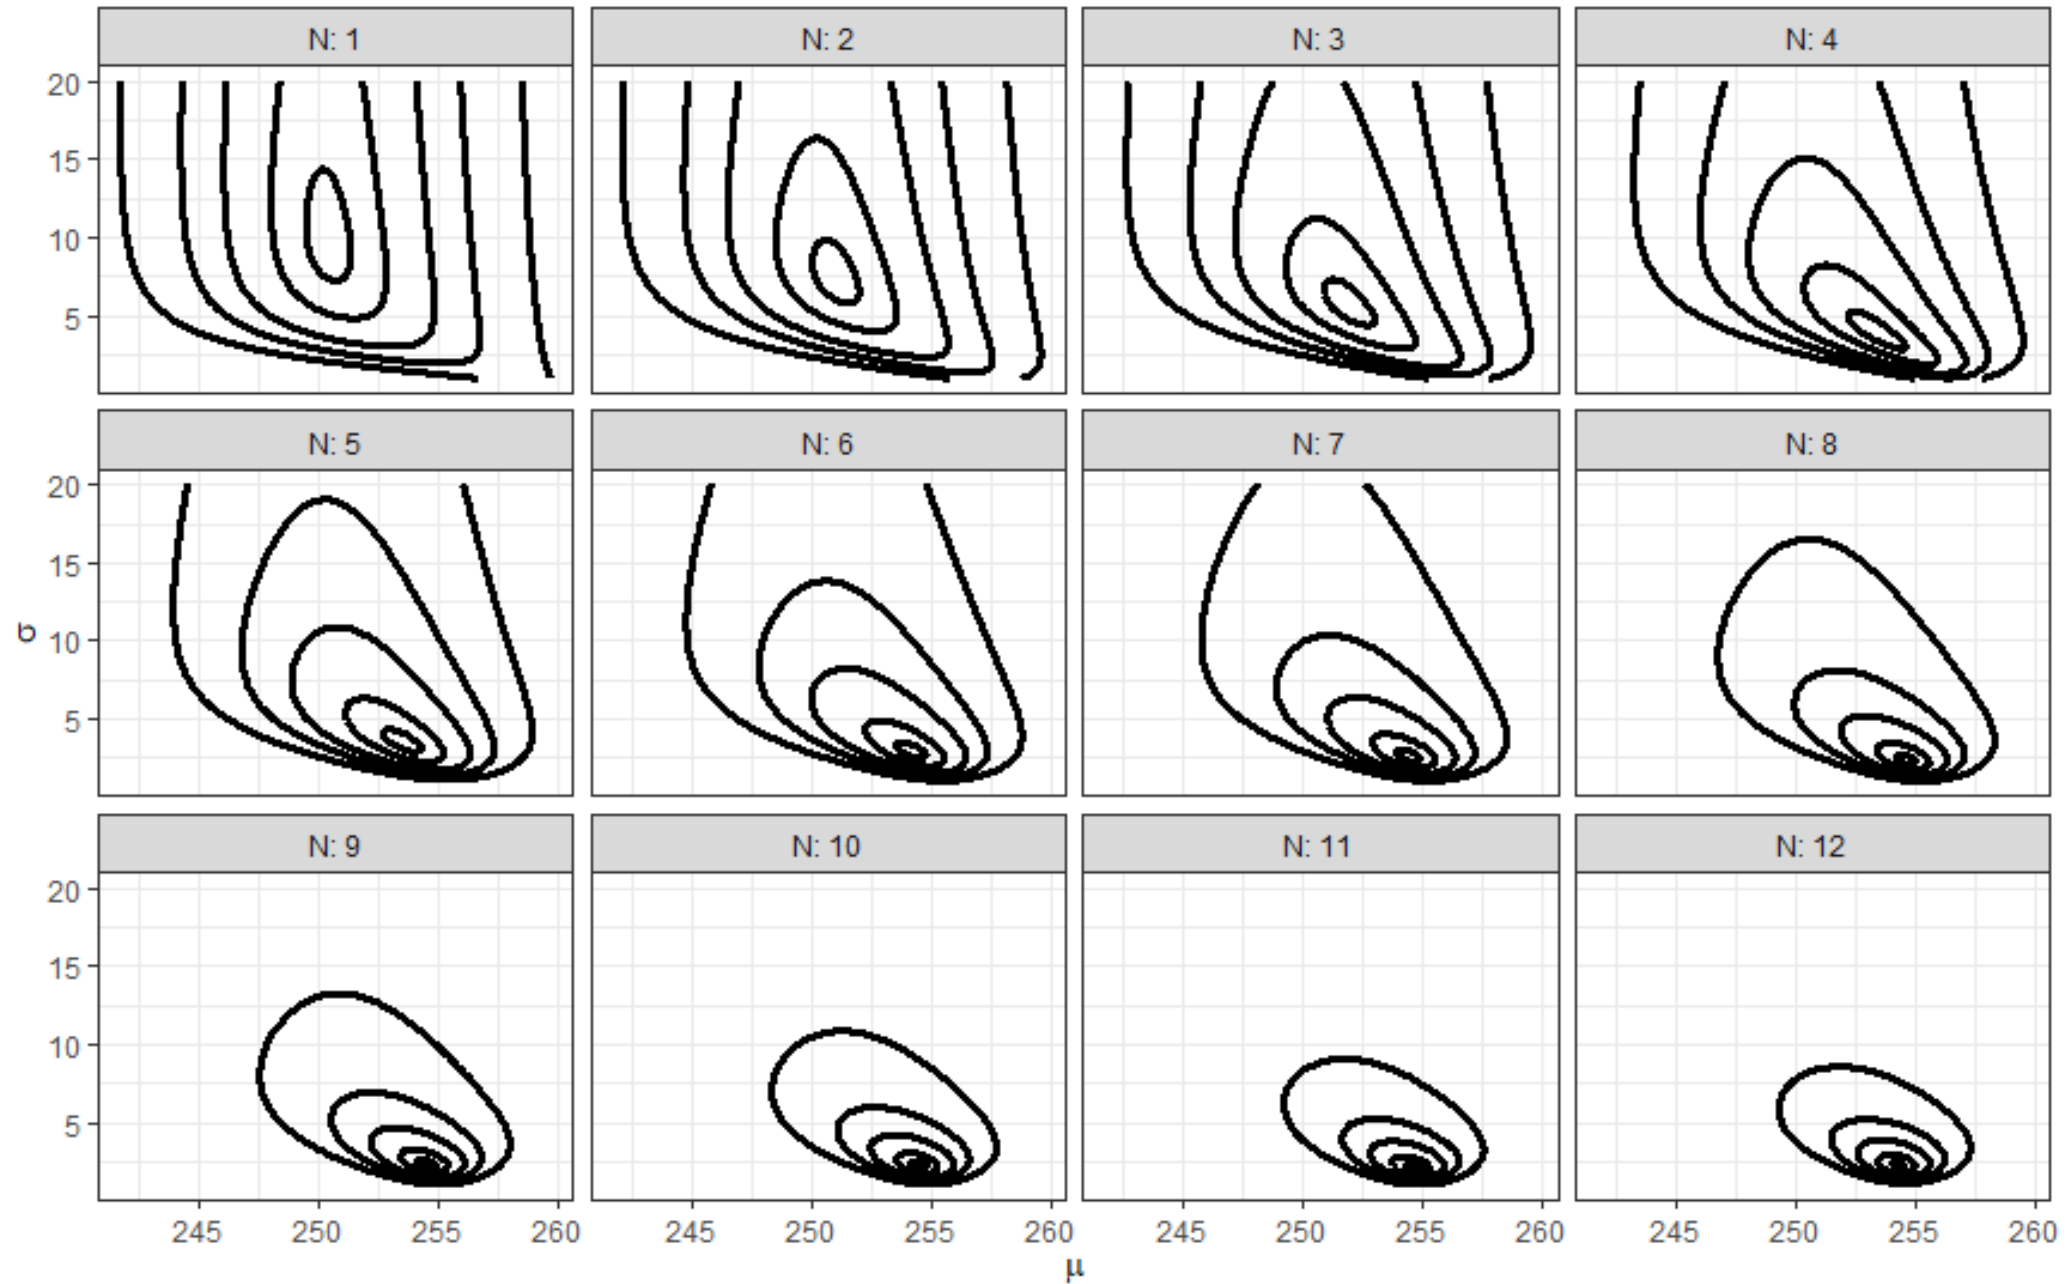
Loop over all parameter pairs with `purrr`

```
77 ### evaluate the log-posterior over the grid  
78 log_post_result <- purrr::map2_dbl(param_grid$mu,  
79                                   param_grid$sigma,  
80                                   eval_logpost,  
81                                   my_info = info_use)  
82
```



Log-posterior surface contour. Grey areas are μ, σ values with posterior probability of less than 0.01%




As a check, the log-posterior surface becomes much more concentrated as the sample size increases.



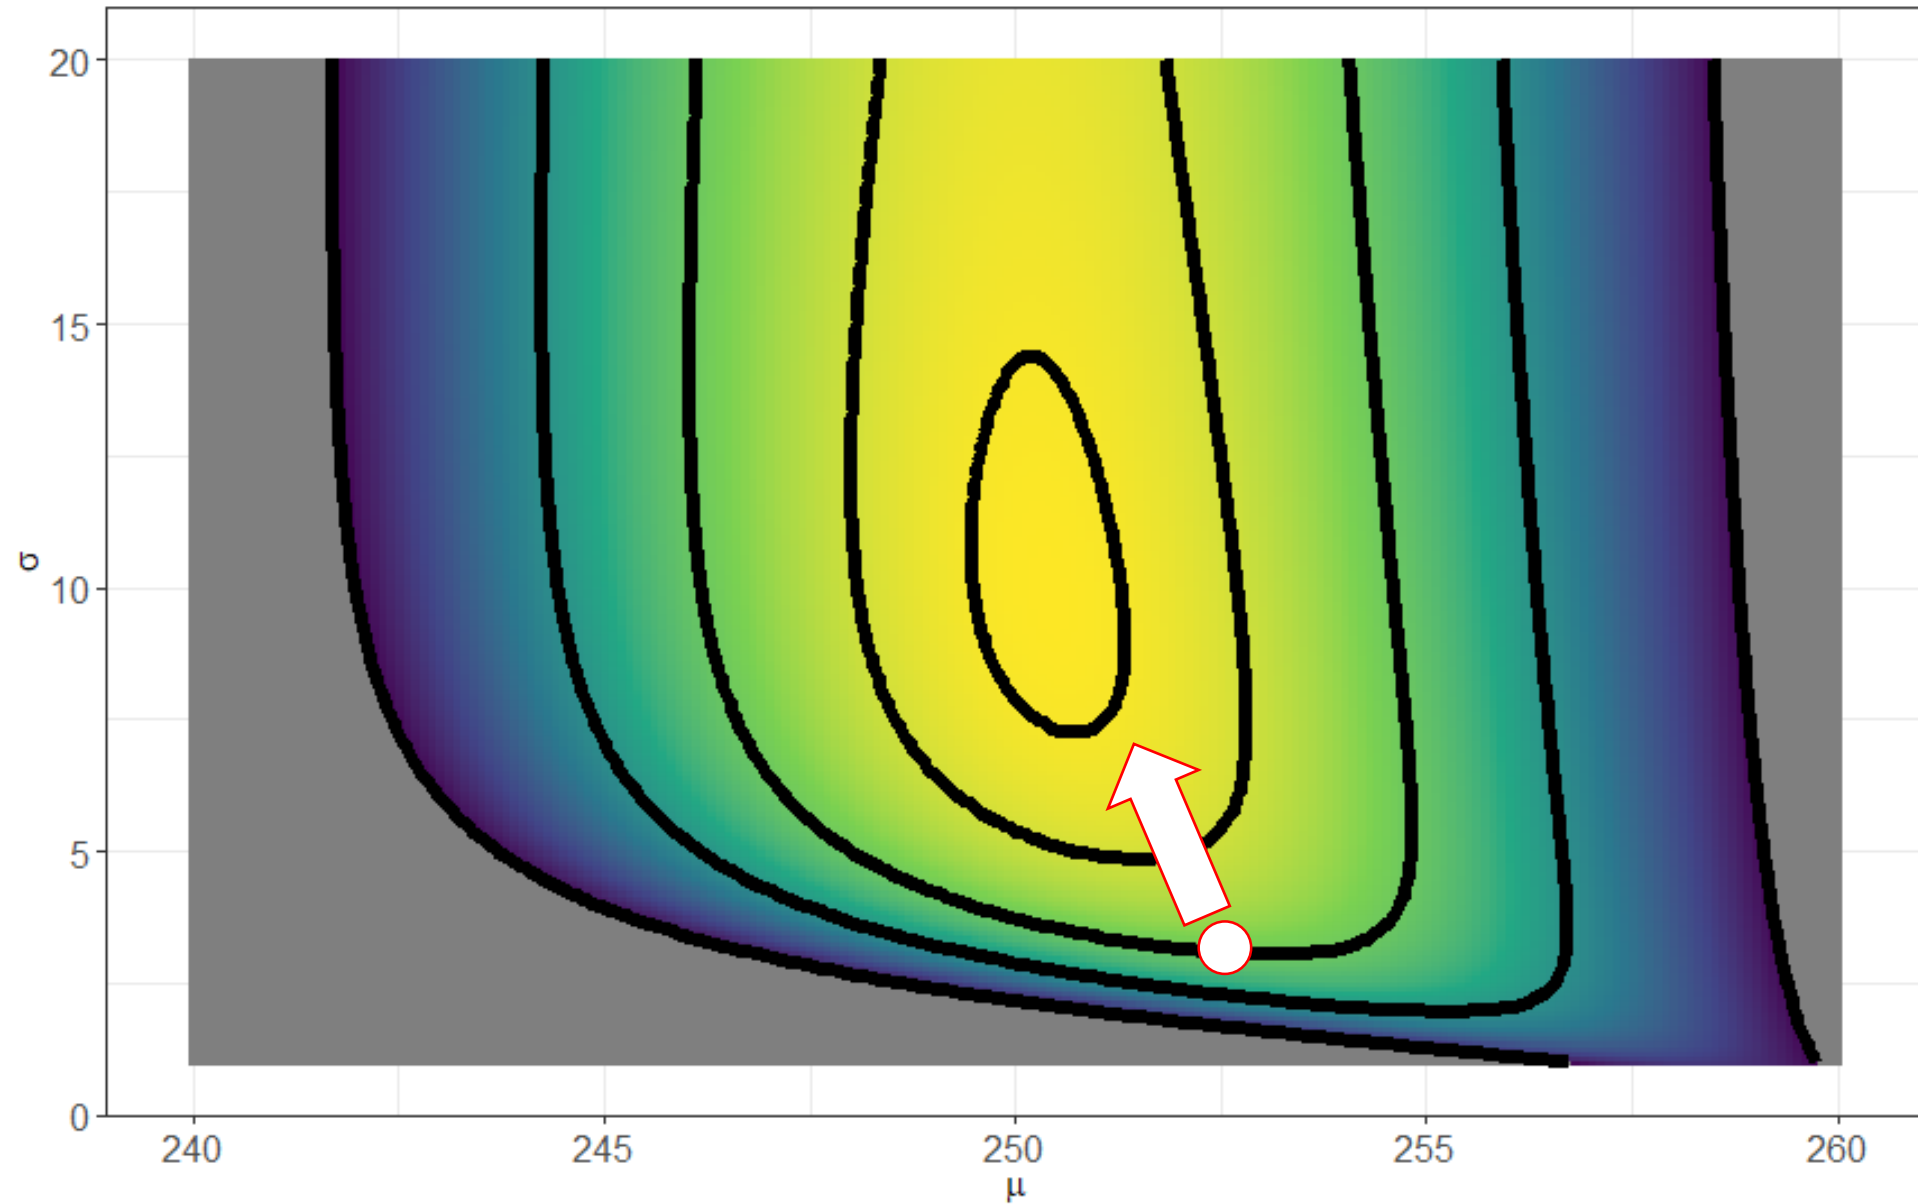
For the $N = 1$ case, how do we find the posterior mode?



Define an initial guess. How should we update our guess?



Steepest ascent! \Rightarrow Find the direction to the peak! \Rightarrow Need to calculate the gradient!



How far should we move in the direction defined by the gradient?

- Numerous algorithms exist for selecting the path or search length.
- Newton-Raphson method sets the search length based on the Hessian!
- For the k -th iteration or step, the $k + 1$ (new) value is:

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \gamma \left[\mathbf{H} \Big|_k \right]^{-1} \mathbf{g} \Big|_k$$

$\gamma \in (0,1)$ is a multiplier
which reduces the step size.

You do NOT need to program the optimization routine...

- We will use the `optim()` function in R to find the posterior mode.
- There are many different optimization algorithms to choose from in R.
- `optim()` is a great starting place for learning about optimization!

`optim()` manages the bookkeeping of the optimization algorithm for us

- Requires that the function we wish to optimize has its **FIRST** argument as a numeric vector.
- `optim()` takes care of modifying the elements of that numeric vector in order to optimize the function of interest.

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```


`optim()` manages the bookkeeping of the optimization algorithm for us

That is why `theta` is the first argument to `my_logpost()` !!!

- Requires that the function we wish to optimize has its **FIRST** argument as a numeric vector.
- `optim()` takes care of modifying the elements of that numeric vector in order to optimize the function of interest.

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try  
217 init_param <- c(info_use$mu_0, 5)  
218  
219 ### find the posterior mode (the MAP) using the optimization scheme  
220 map_result <- optim(init_param,  
221                     my_logpost,  
222                     gr = NULL,  
223                     info_use,  
224                     method = "BFGS",  
225                     hessian = TRUE,  
226                     control = list(fnscale = -1, maxit = 1001))  
227
```

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5) We select an initial guess.
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param, ← That initial guess is the first
221                         my_logpost, argument to optim() .
222                         gr = NULL,
223                         info_use,
224                         method = "BFGS",
225                         hessian = TRUE,
226                         control = list(fnscale = -1, maxit = 1001))
227
```

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost, ← Second argument is the function
222                     gr = NULL,    we wish to optimize.
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
227
```

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
227
```

Third argument is a **function** which returns the gradient vector. If we require the gradient to be evaluated **numerically**, set to `NULL`

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFG",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
227
```

After the `gr` argument, we can pass in **ALL additional inputs** required to evaluate the function we wish to optimize.

Basic syntax of an `optim()` call.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale =
227
```

```
18 my_logpost <- function(theta, my_info)
19 {
20 # the unknown mean is the first parameter
21 lik_mu <- theta[1]
22 # the unknown standard deviation is the second
23 lik_sigma <- theta[2]
24
25 # log-likelihood -> sum up the independent
26 # log-likelihoods
27 log_lik <- sum(dnorm(x = my_info$xobs,
28                     mean = lik_mu,
29                     sd = lik_sigma,
30                     log = TRUE))
31
32 # the log-prior -> sum up the independent priors
33 log_prior <- dnorm(x = lik_mu,
34                    mean = my_info$mu_0,
35                    sd = my_info$tau_0,
36                    log = TRUE) +
37   dunif(x = lik_sigma,
38         min = my_info$sigma_lwr,
39         max = my_info$sigma_upr,
40         log = TRUE)
41
42 # add the log-likelihood and log-prior
43 log_lik + log_prior
44 }
```

Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
227
```


Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
                      my_logpost,
                      gr = NULL,
                      info_use,
                      method = "BFGS",
                      hessian = TRUE,
                      control = list(fnscale = -1, maxit = 1001))
```

Set which optimization method to use. By default `optim()` uses Nelder-Mead. I like to use the Quasi-Newton BFGS algorithm. Try out different methods, see which one you prefer.

Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
```

Tell `optim()` to compute and return the hessian matrix by setting the `hessian` flag to `TRUE`.

Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
```

`control` is a list of parameters which dictate important operating behavior of the algorithm. See `?optim` for a complete list of all available control parameters.

Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
```

`control` is a list of parameters which dictate important operating behavior of the algorithm. See `?optim` for a complete list of all available control parameters.

By default `optim()` seeks to **MINIMIZE** a function, so to tell it to **MAXIMIZE** set the `fnscale` control parameter to `-1`.

Basic syntax of an `optim()` call.

After setting the `gr` argument, all remaining arguments **MUST** be named.

```
216 ### define a starting guess to try
217 init_param <- c(info_use$mu_0, 5)
218
219 ### find the posterior mode (the MAP) using the optimization scheme
220 map_result <- optim(init_param,
221                     my_logpost,
222                     gr = NULL,
223                     info_use,
224                     method = "BFGS",
225                     hessian = TRUE,
226                     control = list(fnscale = -1, maxit = 1001))
```

`control` is a list of parameters which dictate important operating behavior of the algorithm. See `?optim` for a complete list of all available control parameters.

`maxit` controls the maximum number of iterations, default for **derivative** based method is 100.

optim() result

```
> map_result
$par
[1] 250.404541  9.892691

$value
[1] -8.288217

$counts
function gradient
      20      18

$convergence
[1] 0

$message
NULL

$hessian
      [,1]      [,2]
[1,] -0.26021812 -0.02044647
[2,] -0.02044647 -0.02046694
```

To complete the Laplace approximation...

- Define a wrapper which executes the `optim()` call and then calculates the remaining pieces of the Laplace approximation.
- The following code is adapted from the `laplace()` function from the `LearnBayes` package by Jim Albert.
- Great book! [Bayesian Computation with R](#)

Can you decipher what's happening in this function?

```
230 ### define a function for performing the laplace approximation
231 my_laplace <- function(start_guess, logpost_func, ...)
232 {
233   # code adapted from the 'LearnBayes' function 'laplace()'
234   fit <- optim(start_guess,
235               logpost_func,
236               gr = NULL,
237               ...,
238               method = "BFGS",
239               hessian = TRUE,
240               control = list(fnscale = -1, maxit = 1001))
241
242   mode <- fit$par
243   h <- -solve(fit$hessian)
244   p <- length(mode)
245   int <- p/2 * log(2 * pi) + 0.5 * log(det(h)) + logpost_func(mode, ...)
246   list(mode = mode,
247        var_matrix = h,
248        log_evidence = int,
249        converge = ifelse(fit$convergence == 0,
250                          "YES",
251                          "NO"),
252        iter_counts = fit$counts[1])
253 }
```


Can you decipher what's happening in this function?

```
230 ### define a function for performing the laplace approximation
231 my_laplace <- function(start_guess, logpost_func, ...)
232 {
233   # code adapted from the 'LearnBayes' function 'laplace()'
234   fit <- optim(start_guess,
235               logpost_func,
236               gr = NULL,
237               ...,
238               method = "BFGS",
239               hessian = TRUE,
240
241   mode <- f
242   h <- -sol
243   p <- length(mode)
244   int <- p/2 * log(2 * pi) + 0.5 * log(det(h)) + logpost_func(mode, ...)
245   list(mode = mode,
246        var_matrix = h,
247        log_evidence = int,
248        converge = ifelse(fit$convergence == 0,
249                          "YES",
250                          "NO"),
251        iter_counts = fit$counts[1])
252 }
253 }
```

We will return to this aspect of the Laplace approximation in a future lecture!

Perform the Laplace approximation

```
> laplace_result_N01 <- my_laplace(init_param, my_logpost, info_use)
> laplace_result_N01
$mode
[1] 250.404541    9.892691

$var_matrix
      [,1]      [,2]
[1,]  4.170279 -4.166109
[2,] -4.166109 53.021223

$log_evidence
[1] -3.791876

$converge
[1] "YES"

$iter_counts
function
      20
```

WHAT IS THIS???????

```
> laplace_result_N01 <- my_laplace(init_param, my_logpost, info_use)
> laplace_result_N01
$mode
[1] 250.404541    9.892691


$var_matrix
      [,1]      [,2]
[1,]  4.170279 -4.166109
[2,] -4.166109 53.021223

$log_evidence
[1] -3.791876


$converge
[1] "YES"

$iter_counts
function
      20
```


Approximate log-posterior based on the MVN Laplace approximation.




Compare true log-posterior surface (black) with the MVN approximate log-posterior (dashed red)



Compare true log-posterior surface (black) with the MVN approximate log-posterior (dashed red)



Compare true log-posterior surface (black) with the MVN approximate log-posterior (dashed red)



Confirm similarities and differences by drawing samples and visualizing results

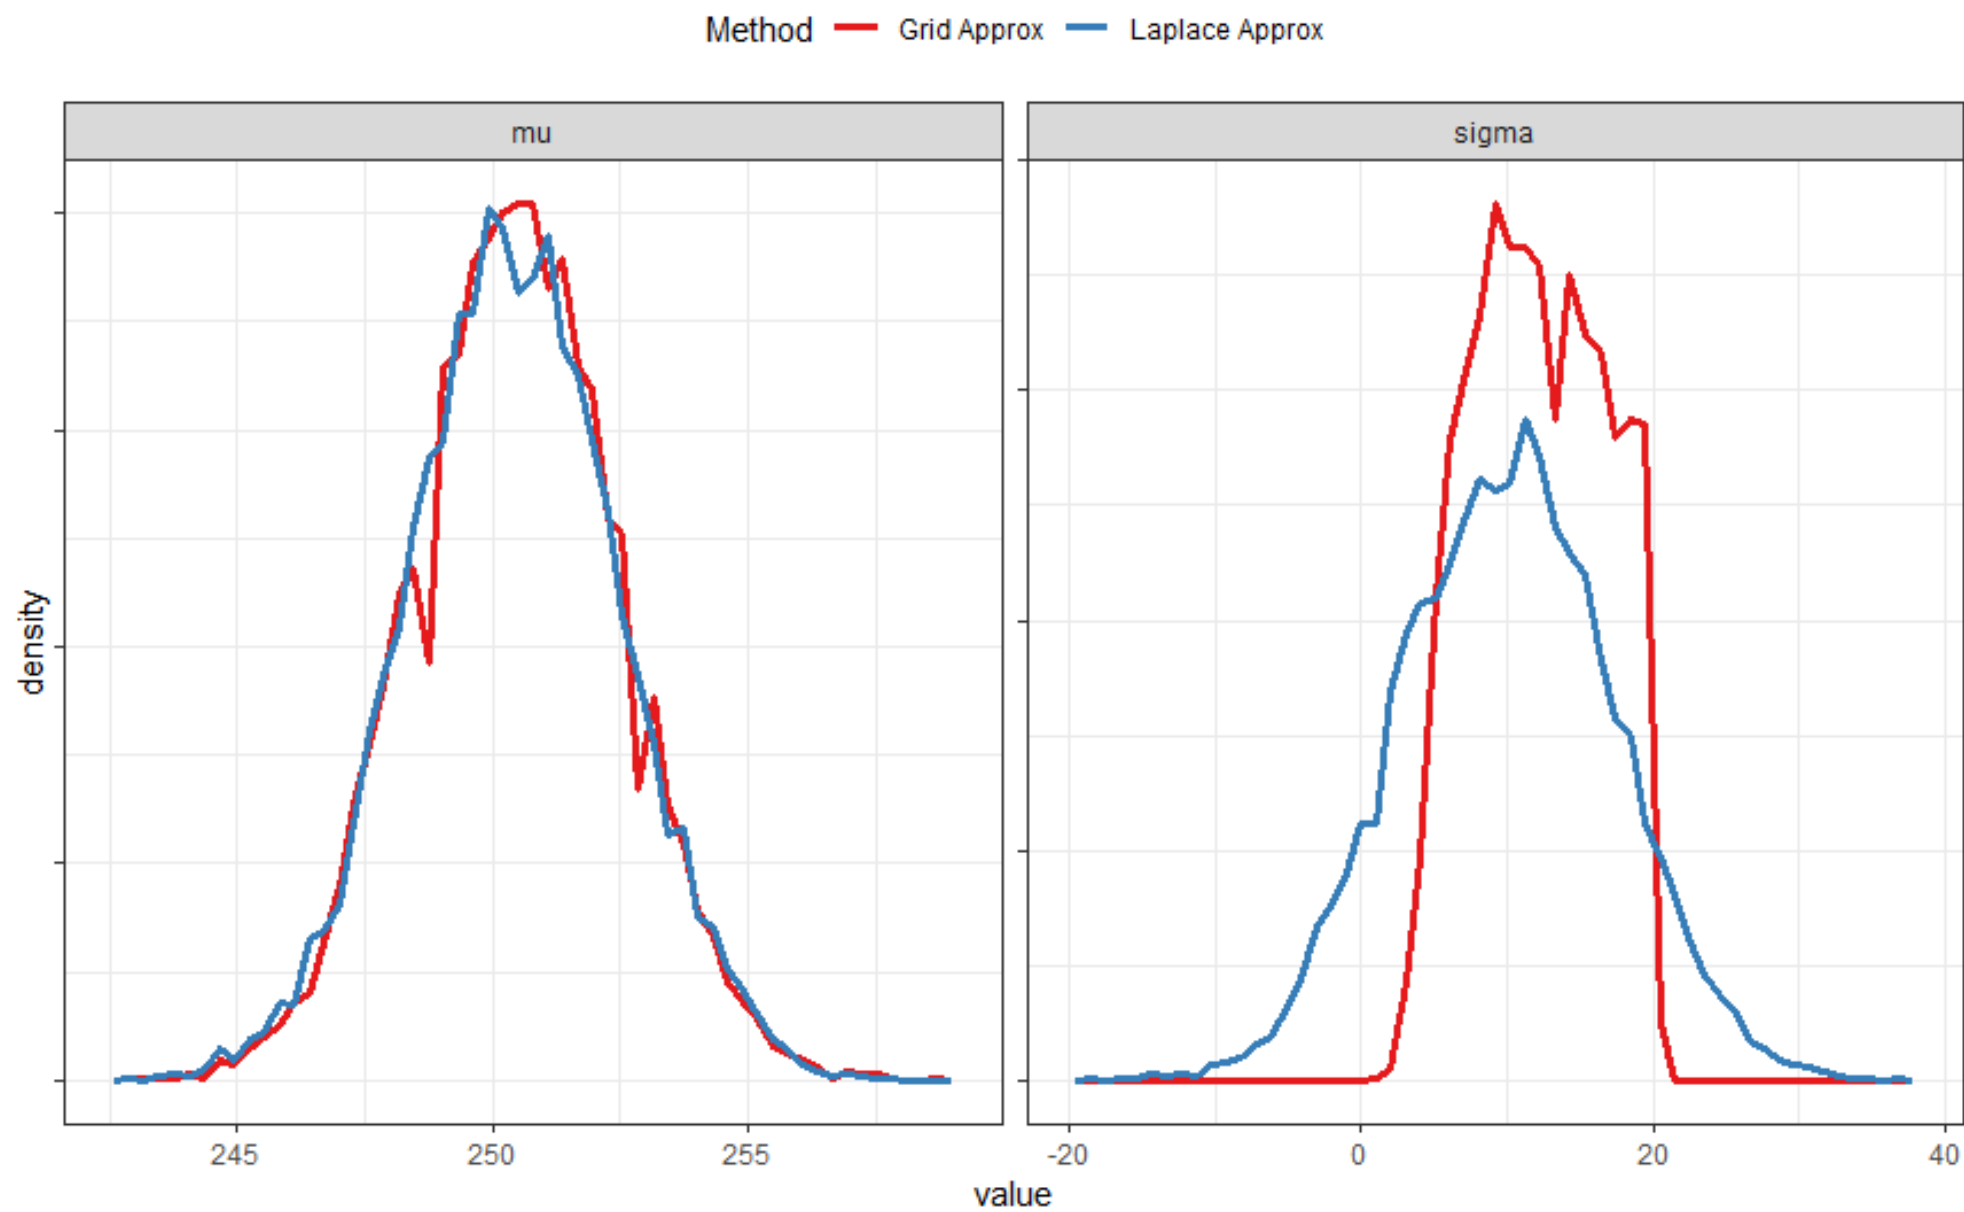
```
348 ### random draws from the approximate MVN posterior
349 set.seed(5002)
350 post_mvn_samples <- MASS::mvrnorm(n = 1e4,
351                                   mu = laplace_result_N01$mode,
352                                   Sigma = laplace_result_N01$var_matrix) %>%
353   as.data.frame() %>% tbl_df() %>%
354   purrr::set_names(c("mu", "sigma"))
355
356 post_mvn_samples
357
358 ### draw samples from the grid approximation directly to compare to the
359 ### approximate MVN posterior samples
360
361 grid_approx_result <- param_grid %>%
362   mutate(log_post = log_post_result) %>%
363   mutate(log_post_2 = log_post - max(log_post))
364
365 set.seed(5003)
366 direct_sample_id <- sample(1:nrow(param_grid),
367                             size = 1e4,
368                             replace = TRUE,
369                             prob = exp(grid_approx_result$log_post_2))
370
371 grid_approx_samples <- grid_approx_result %>%
372   slice(direct_sample_id)
373
```

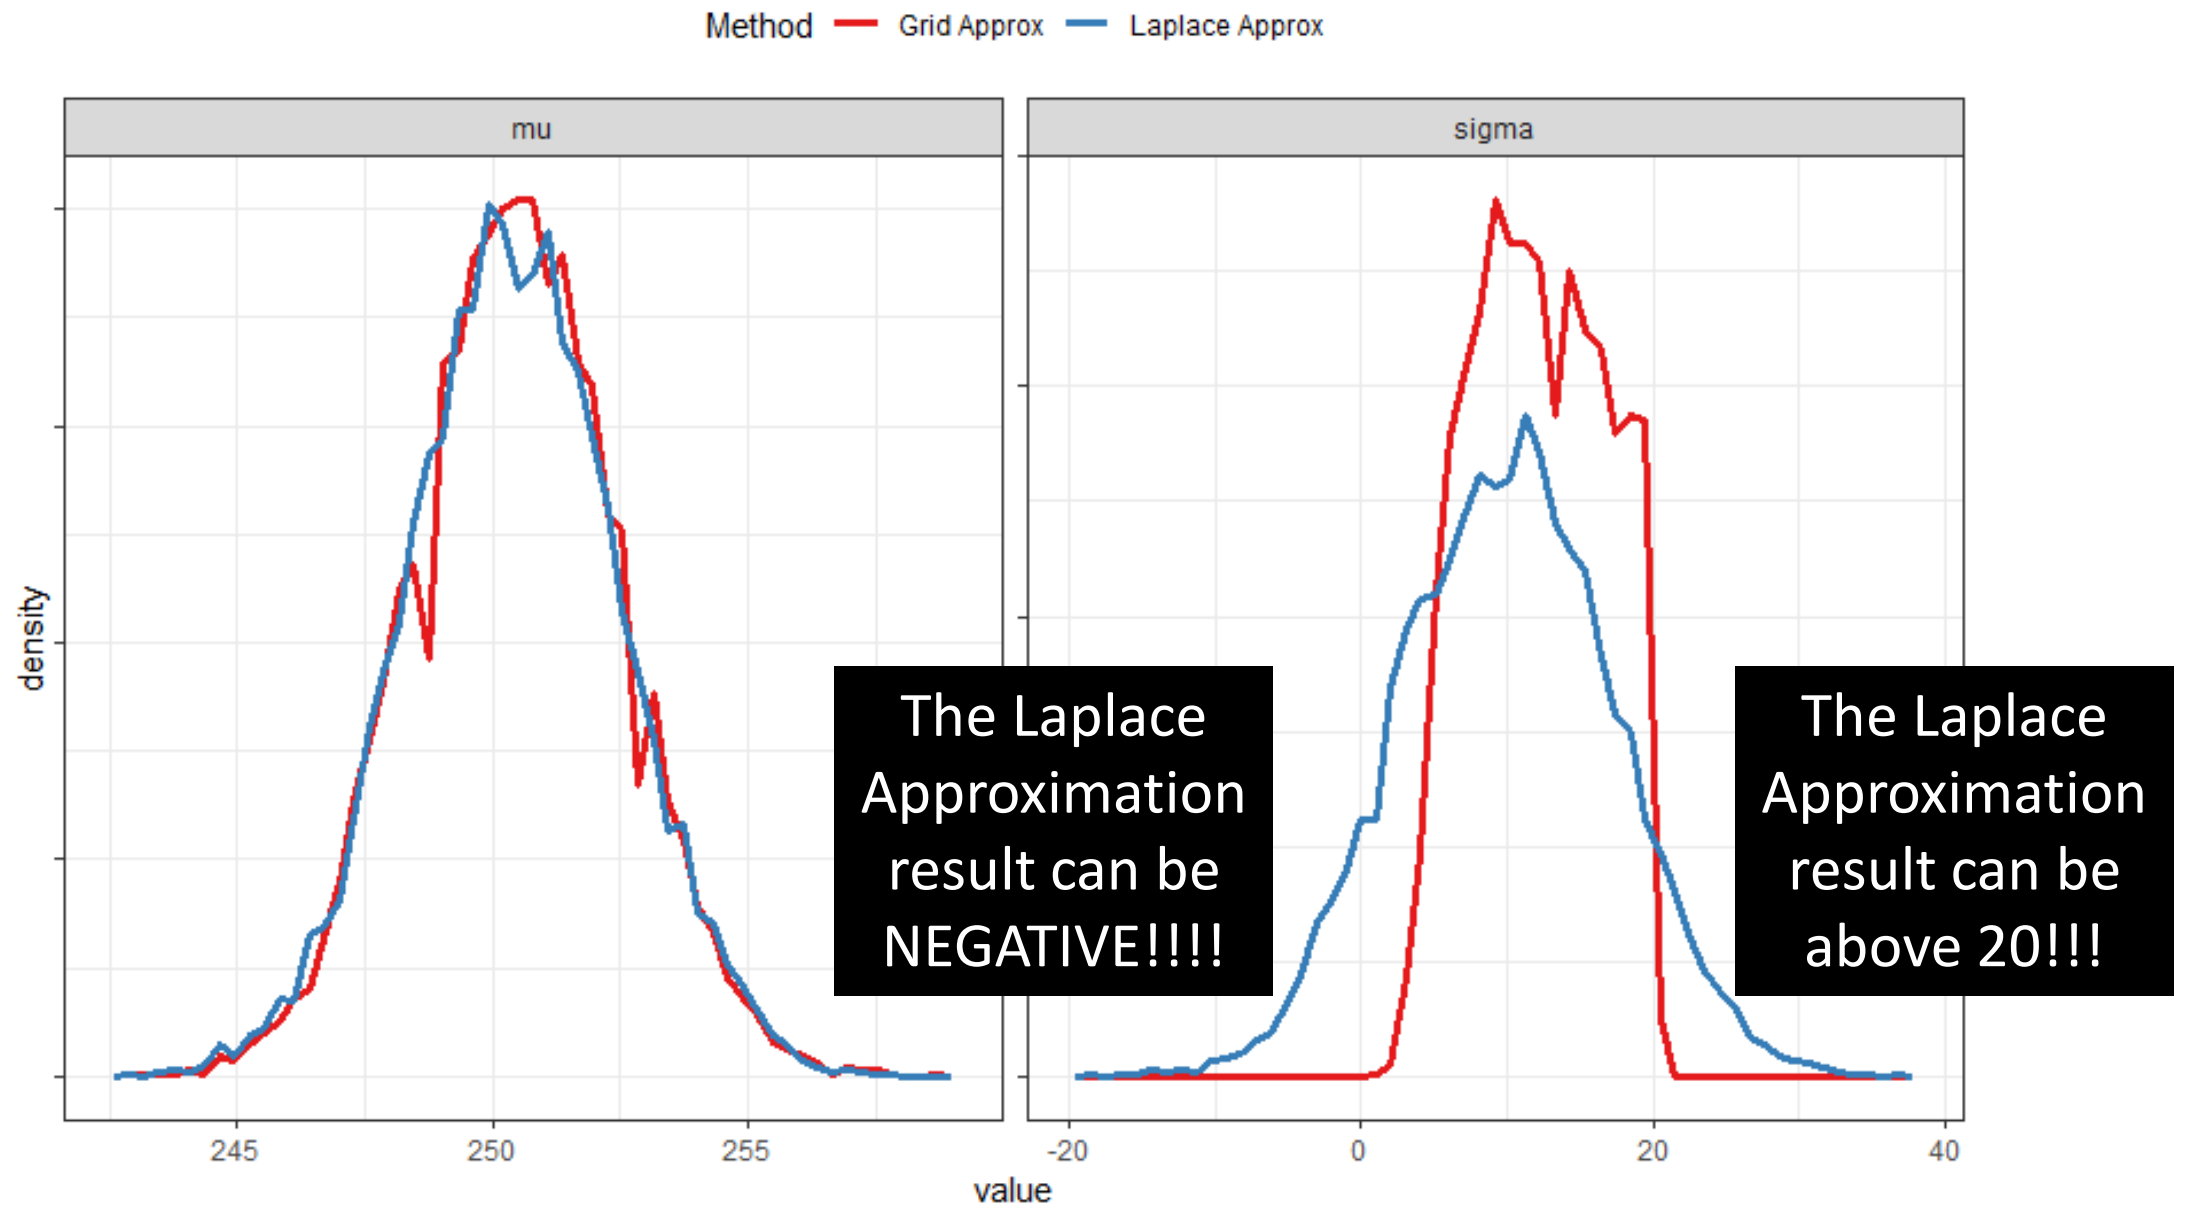

Confirm similarities and differences by drawing samples and visualizing results

```
348 ### random draws from the approximate MVN posterior
349 set.seed(5002)
350 post_mvn_samples <- MASS::mvrnorm(n = 1e4,
351                                   mu = laplace_result_N01$mode,
352                                   Sigma = laplace_result_N01$var_matrix) %>%
353   as.data.frame() %>% tbl_df() %>%
354   purrr::set_names(c("mu", "sigma"))
355
356 post_mvn_samples
357
358 ### draw samples from the grid approximation directly to compare to the
359 ### approximate MVN posterior samples
360
361 grid_approx_result <- param_grid %>%
362   mutate(log_post = log_post_result) %>%
363   mutate(log_post_2 = log_post - max(log_post))
364
365 set.seed(5003)
366 direct_sample_id <- sample(1:nrow(param_grid),
367                             size = 1e4,
368                             replace = TRUE,
369                             prob = exp(grid_approx_result$log_post_2))
370
371 grid_approx_samples <- grid_approx_result %>%
372   slice(direct_sample_id)
373
```

Samples based on the approximate MVN distribution directly.

Sample with replacement with row weights equal to the posterior probability of that row.





What's going on?

- Each variable in a Multivariate Normal distribution has a Gaussian distribution.
- Gaussian variables are unbounded: $-\infty \rightarrow +\infty$ are allowed!
- The natural lower 0 bound on σ is therefore ignored, as is the upper bound imposed by the UNIFORM prior.

How can we have the imposed UNIFORM constraint satisfied?

- Apply a transformation!
- Use the change-of-variables procedure to transform from the bounded σ to a new unbounded variable ϕ .

$$\sigma = g^{-1}(\phi), \phi = g(\sigma)$$

- Apply the Laplace approximation to the joint posterior distribution on μ, ϕ since both are unbounded.

With the change-of-variables procedure we can still use our original prior on σ

The joint-posterior on μ, ϕ can therefore be written based on the joint posterior μ, σ :

$$p(\mu, \phi | \mathbf{x}) = \prod_{n=1}^N \{p(x_n | \mu, g^{-1}(\phi))\} p(\mu | \mu_0, \tau_0) p(g^{-1}(\phi) | l, u) \cdot \left| \frac{d}{d\phi} (g^{-1}(\phi)) \right|$$

With the change-of-variables procedure we can still use our original prior on σ

The joint-posterior on μ, ϕ can therefore be written based on the joint posterior μ, σ :

$$p(\mu, \phi | \mathbf{x}) = \underbrace{\prod_{n=1}^N \{p(x_n | \mu, g^{-1}(\phi))\} p(\mu | \mu_0, \tau_0) p(g^{-1}(\phi) | l, u)}_{\text{"Back-transform" } \phi \text{ to } \sigma \text{ and substitute into the joint-posterior just as before when we were working with } \sigma \text{ directly.}} \cdot \left| \frac{d}{d\phi} (g^{-1}(\phi)) \right|$$

“Back-transform” ϕ to σ and substitute into the joint-posterior just as before when we were working with σ directly.

With the change-of-variables procedure we can still use our original prior on σ

The joint-posterior on μ, ϕ can therefore be written based on the joint posterior μ, σ :

$$p(\mu, \phi | \mathbf{x}) = \prod_{n=1}^N \{p(x_n | \mu, g^{-1}(\phi))\} p(\mu | \mu_0, \tau_0) p(g^{-1}(\phi) | l, u) \cdot \left| \frac{d}{d\phi} (g^{-1}(\phi)) \right|$$

Derivative of the
transformation
with respect to ϕ

Use a logit-transformation to satisfy the lower and upper bounds on σ

- The logit, or log-odds transformation maps a lower and upper bounded variable to an unbounded variable.

$$\phi = \text{logit}(\psi) = \log \left[\frac{\psi}{1 - \psi} \right]$$

$$\psi = \text{logit}^{-1}(\phi) = \text{logistic}(\phi) = \frac{1}{1 + \exp(-\phi)}$$

- **We will discuss the logit-transformation in more detail when we talk about classification.**

Use a logit-transformation to satisfy the lower and upper bounds on σ

Transformation

$$\phi = \text{logit} \left(\frac{\sigma - l}{u - l} \right) = g(\sigma)$$

Inverse transformation

$$\sigma = l + (u - l) \cdot \text{logit}^{-1}(\phi) = g^{-1}(\phi)$$

Derivative

$$\frac{d\sigma}{d\phi} = (u - l) \cdot \text{logit}^{-1}(\phi) \cdot (1 - \text{logit}^{-1}(\phi))$$

Define a new function `my_logpost_cv()` which calculates the log-posterior on μ, ϕ

```
399 my_logpost_cv <- function(phi, my_info)
400 {
401   # the unknown mean is the first parameter
402   lik_mu <- phi[1]
403   # the unknown logit-transformed standard deviation
404   # is the second, back transform to sigma
405   lik_sigma <- my_info$sigma_lwr +
406     (my_info$sigma_upr - my_info$sigma_lwr) * boot::inv.logit(phi[2])
407
408   # log-likelihood -> sum up the independent
409   # log-likelihoods
410   log_lik <- sum(dnorm(x = my_info$xobs,
411                       mean = lik_mu,
412                       sd = lik_sigma,
413                       log = TRUE))
414
415   # the log-prior -> sum up the independent priors
416   log_prior <- dnorm(x = lik_mu,
417                     mean = my_info$mu_0,
418                     sd = my_info$tau_0,
419                     log = TRUE) +
420     dunif(x = lik_sigma,
421          min = my_info$sigma_lwr,
422          max = my_info$sigma_upr,
423          log = TRUE)
424
425   # add the log-likelihood and log-prior and account
426   # for the derivative adjustment
427   deriv_adjust <- log(my_info$sigma_upr - my_info$sigma_lwr) +
428     log(boot::inv.logit(phi[2])) +
429     log(1 - boot::inv.logit(phi[2]))
430
431   log_lik + log_prior + deriv_adjust
432 }
```


The log-likelihood and log-priors are evaluated just as they were in the original function!

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```

```
399 my_logpost_cv <- function(phi, my_info)
400 {
401   # the unknown mean is the first parameter
402   lik_mu <- phi[1]
403   # the unknown logit-transformed standard deviation
404   # is the second, back transform to sigma
405   lik_sigma <- my_info$sigma_lwr +
406     (my_info$sigma_upr - my_info$sigma_lwr) * boot::inv.logit(phi[2])
407
408   # log-likelihood -> sum up the independent
409   # log-likelihoods
410   log_lik <- sum(dnorm(x = my_info$xobs,
411                       mean = lik_mu,
412                       sd = lik_sigma,
413                       log = TRUE))
414
415   # the log-prior -> sum up the independent priors
416   log_prior <- dnorm(x = lik_mu,
417                     mean = my_info$mu_0,
418                     sd = my_info$tau_0,
419                     log = TRUE) +
420     dunif(x = lik_sigma,
421          min = my_info$sigma_lwr,
422          max = my_info$sigma_upr,
423          log = TRUE)
424
425   # add the log-likelihood and log-prior and account
426   # for the derivative adjustment
427   deriv_adjust <- log(my_info$sigma_upr - my_info$sigma_lwr) +
428     log(boot::inv.logit(phi[2])) +
429     log(1 - boot::inv.logit(phi[2]))
430
431   log_lik + log_prior + deriv_adjust
432 }
```

The new function “back-transforms” ϕ to σ


```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                        mean = lik_mu,
29                        sd = lik_sigma,
30                        log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```



```
399 my_logpost_cv <- function(phi, my_info)
400 {
401   # the unknown mean is the first parameter
402   lik_mu <- phi[1]
403   # the unknown logit-transformed standard deviation
404   # is the second, back transform to sigma
405   lik_sigma <- my_info$sigma_lwr +
406     (my_info$sigma_upr - my_info$sigma_lwr) * boot::inv.logit(phi[2])
407
408   # log-likelihood -> sum up the independent
409   # log-likelihoods
410   log_lik <- sum(dnorm(x = my_info$xobs,
411                       mean = lik_mu,
412                       sd = lik_sigma,
413                       log = TRUE))
414
415   # the log-prior -> sum up the independent priors
416   log_prior <- dnorm(x = lik_mu,
417                     mean = my_info$mu_0,
418                     sd = my_info$tau_0,
419                     log = TRUE) +
420     dunif(x = lik_sigma,
421          min = my_info$sigma_lwr,
422          max = my_info$sigma_upr,
423          log = TRUE)
424
425   # add the log-likelihood and log-prior and account
426   # for the derivative adjustment
427   deriv_adjust <- log(my_info$sigma_upr - my_info$sigma_lwr) +
428     log(boot::inv.logit(phi[2])) +
429     log(1 - boot::inv.logit(phi[2]))
430
431   log_lik + log_prior + deriv_adjust
432 }
```

The new function calculates the log-derivative of σ with respect to ϕ

```
18 my_logpost <- function(theta, my_info)
19 {
20   # the unknown mean is the first parameter
21   lik_mu <- theta[1]
22   # the unknown standard deviation is the second
23   lik_sigma <- theta[2]
24
25   # log-likelihood -> sum up the independent
26   # log-likelihoods
27   log_lik <- sum(dnorm(x = my_info$xobs,
28                       mean = lik_mu,
29                       sd = lik_sigma,
30                       log = TRUE))
31
32   # the log-prior -> sum up the independent priors
33   log_prior <- dnorm(x = lik_mu,
34                     mean = my_info$mu_0,
35                     sd = my_info$tau_0,
36                     log = TRUE) +
37     dunif(x = lik_sigma,
38          min = my_info$sigma_lwr,
39          max = my_info$sigma_upr,
40          log = TRUE)
41
42   # add the log-likelihood and log-prior
43   log_lik + log_prior
44 }
```




```
399 my_logpost_cv <- function(phi, my_info)
400 {
401   # the unknown mean is the first parameter
402   lik_mu <- phi[1]
403   # the unknown logit-transformed standard deviation
404   # is the second, back transform to sigma
405   lik_sigma <- my_info$sigma_lwr +
406     (my_info$sigma_upr - my_info$sigma_lwr) * boot::inv.logit(phi[2])
407
408   # log-likelihood -> sum up the independent
409   # log-likelihoods
410   log_lik <- sum(dnorm(x = my_info$xobs,
411                       mean = lik_mu,
412                       sd = lik_sigma,
413                       log = TRUE))
414
415   # the log-prior -> sum up the independent priors
416   log_prior <- dnorm(x = lik_mu,
417                     mean = my_info$mu_0,
418                     sd = my_info$tau_0,
419                     log = TRUE) +
420     dunif(x = lik_sigma,
421          min = my_info$sigma_lwr,
422          max = my_info$sigma_upr,
423          log = TRUE)
424
425   # add the log-likelihood and log-prior and account
426   # for the derivative adjustment
427   deriv_adjust <- log(my_info$sigma_upr - my_info$sigma_lwr) +
428     log(boot::inv.logit(phi[2])) +
429     log(1 - boot::inv.logit(phi[2]))
430
431   log_lik + log_prior + deriv_adjust
432 }
```

Let's first visualize the log-posterior for the unbounded parameters μ, ϕ

```
465 ### define a new grid within the unbounded parameter  
466 ### space, and evaluate the (mu,phi) log-posterior on this  
467 ### new grid  
468 phi_grid <- expand.grid(mu = seq(240, 260, length.out = 201),  
469                          phi = seq(boot::logit(1e-3),  
470                                   boot::logit(0.999),  
471                                   length.out = 201),  
472                          KEEP.OUT.ATTRS = FALSE,  
473                          stringsAsFactors = FALSE) %>%  
474   as.data.frame() %>% tbl_df()  
475
```

Log-posterior for the unbounded parameters μ, ϕ



Perform the Laplace approximation within the unbounded domain

```
> cv_laplace_result_N01 <- my_laplace(c(250, log(5)), my_logpost_cv, info_use)
> cv_laplace_result_N01
$mode
[1] 250.37888274 -0.05571467

$var_matrix
      [,1]      [,2]
[1,] 3.9857940 -0.3926613
[2,] -0.3926613 1.1601275

$log_evidence
[1] -4.14531

$converge
[1] "YES"

$iter_counts
function
      18
```

Perform the Laplace approximation within the unbounded domain

```
> cv_laplace_result_N01 <- my_laplace(c(250, log(5)), my_logpost_cv, info_use)
> cv_laplace_result_N01
$mode
[1] 250.37888274 -0.05571467


$var_matrix
      [,1] [,2]
[1,] 3.9857940 -0.3926613
[2,] -0.3926613 1.1601275

$log_evidence
[1] -4.14531

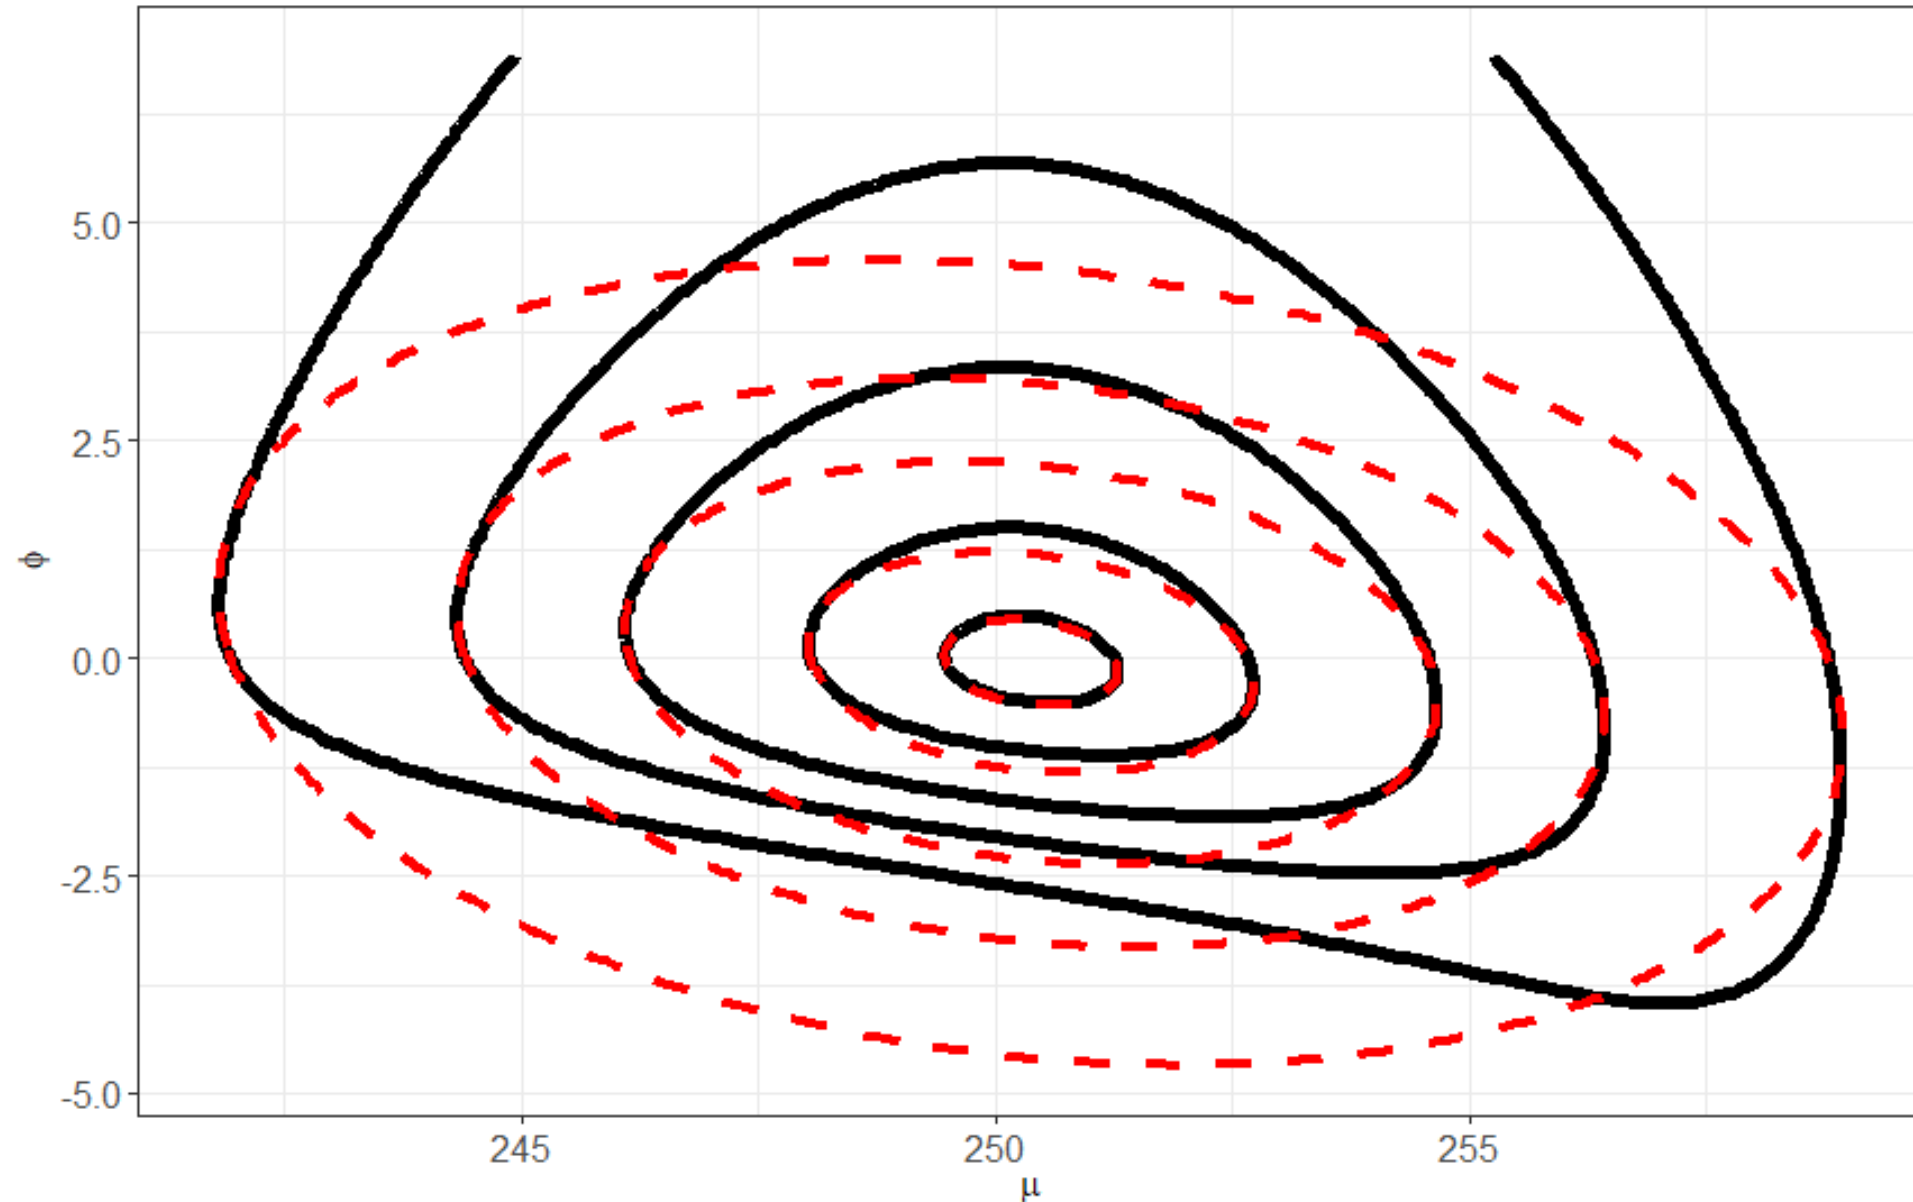
$converge
[1] "YES"

$iter_counts
function
      18
```


Approximate log-posterior based on the MVN Laplace approximation in the unbounded parameter space.



Compare true log-posterior surface (black) with the MVN approximate log-posterior (dashed red)



Compare true log-posterior surface (black) with the MVN approximate log-posterior (dashed red)



Generate random samples from the approximate MVN posterior on μ, ϕ

```
560 ### draw posterior samples from the MVN approximate posterior  
561 ### on the unbounded variables  
562 set.seed(5004)  
563  
564 post_cv_mvn_samples <- MASS::mvrnorm(n = 1e4,  
565                                     mu = cv_laplace_result_N01$mode,  
566                                     sigma = cv_laplace_result_N01$var_matrix) %>%  
567   as.data.frame() %>% tbl_df() %>%  
568   purrr::set_names(c("mu", "logit_sigma"))  
569
```

“Back-transform” the posterior ϕ samples to σ and compare with the previous Laplace approximation and “true” grid approximation results

```
575 grid_approx_samples %>%
576   select(mu, sigma) %>%
577   mutate(type = "Grid Approx") %>%
578   bind_rows(post_mvn_samples %>%
579     mutate(type = "Laplace Approx")) %>%
580   bind_rows(post_cv_mvn_samples %>%
581     mutate(sigma = info_use$sigma_lwr +
582       (info_use$sigma_upr - info_use$sigma_lwr) * boot::inv.logit(logit_sigma)) %>%
583     select(mu, sigma) %>%
584     mutate(type = "Laplace Approx with transformation")) %>%
585   tibble::rowid_to_column("post_id") %>%
586   tidyr::gather(key = "key", value = "value", -post_id, -type) %>%
587   ggplot(mapping = aes(x = value)) +
588   geom_freqpoly(mapping = aes(group = interaction(key, type),
589     color = type,
590     y = stat(density)),
591     bins = 55, size = 1.15) +
592   facet_wrap(~ key, scales = "free") +
593   ggthemes::scale_color_colorblind("Method") +
594   theme_bw() +
595   theme(legend.position = "top",
596     axis.text.y = element_blank())
597
```

